

# **G22.2250-001**

## **Operating Systems**

Computer and Operating System Structures  
Processes, Threads, and Process Cooperation

September 11, 2007

### Outline

---


#### Announcements

- TA: Jay Chen (jchen@cs), 7<sup>th</sup> Floor 715 B'way
  - Office hours: Mondays 5:00 – 6:00pm, Thursdays 4:00 – 5:00pm
  - Primary point of contact for questions about Nachos labs
- Computer system structures
  - (Review) I/O, Storage, Protection
- Operating system structures
- Processes
  - process scheduling
- Threads
  - Multithreading models
- Process cooperation
  - Shared memory vs. message passing

[ Silberschatz/Galvin/Gagne: Chapters 1, 2, 3.1-3.3, 4.1-4.5, 6.1 ]

(Review)

## The Hardware of a Modern Computer System



9/11/2007

3

(Review)

## Computer-System Structures (1): Input/Output

- CPU interaction with I/O devices via device controllers
  - Special-purpose processors with local storage and registers
  - CPU requests service by writing to control registers
  - Device controllers **interrupt** the CPU upon completion of action
    - CPU support for interrupts
      - Saving state of currently running process (registers)
      - Vectoring to interrupt service routine (ISR)
      - Return from ISR restores process state, resuming as if the interrupt never happened
    - **Traps** are like interrupts, except triggered by special CPU instructions
- I/O operations can be synchronous or asynchronous
  - **Direct memory access** (DMA) mechanisms help reduce CPU involvement
- **Memory-mapped I/O** permits controller storage to be accessed using CPU ld/st instructions
  - Essential for high-speed I/O devices (e.g., video)

9/11/2007

4

## Computer-System Structures (2): Storage

---

- Primary storage: *Main memory* (volatile)
  - accessed directly using load/store instructions
    - 1 cycle (registers), 2-5 cycles (cache), 20-50 cycles (RAM)
    - *before*: only one outstanding memory operation, CPU waits for completion
    - *now*: several outstanding operations
- Secondary storage: *Disks* (non-volatile)
  - accessed using a disk controller
  - supports random access but with non-uniform cost
- Tertiary storage: *Tapes, Optical disks* (non-volatile)
  - typically used only for backup
  - very inefficient support for random access
- Organized as a hierarchy
  - small amount of faster, more expensive storage closer to the CPU
  - larger amounts of slower, less expensive storage further away

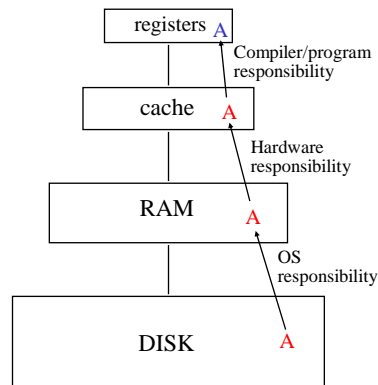
9/11/2007

5

## Storage Hierarchy

---

- Rationale
  - keep CPU busy: lots of fast memory
  - keep system cost down
- How does it work
  - *caching*: upon access, move datum or instruction and its neighbors into higher levels of the hierarchy
  - *replacement* when a level fills up
  - copies need to be kept **coherent**
- Why does it work
  - Real programs demonstrate *locality*
    - e.g.: rows and columns of a matrix
    - e.g.: sequential instructions
  - once a *datum* or *instruction* is used, things “near” them are likely to be used “soon”



9/11/2007

6

## Computer-System Structures (3): Protection

---

- Goal: Prevent user processes from accidentally/maliciously damaging
  - the OS structures
  - parts of other process's memory space
  - other user's I/O devices
- Mechanisms address different ways in which protection breaks down
  1. **dual-mode operation**
    - Prevent user process taking over part of the OS and using this to overwrite other processes or even modify the OS itself (as in MS-DOS)
  2. **privileged instructions**
    - Prevent user process intervening in I/O of another process via control of the I/O handlers and indirectly causing damage
  3. **memory protection**
    - Prevent user process directly accessing another user process' storage
  4. **CPU protection via timers**
    - Prevent hanging the OS -- e.g., via an infinite loop


9/11/2007

7

## Protection Mechanisms (1): Dual-mode Operation and Privileged Instructions

---

- Dual-mode operation
  - *supervisor* and *user* modes
  - system starts off in supervisor mode and reenters it for interrupt processing
  - operating system gains control in supervisor mode
- Privileged instructions
  - restrict use of certain instructions to supervisor mode
    - I/O, including interrupt control
      - exception is instructions which generate interrupts
      - may be done by memory mapping
    - affect memory mapping
    - affect CPU mode (user/supervisor)
  - hardware support crucial for performance and for atomicity




9/11/2007

8

## Protection Mechanisms (2): Memory Protection

---

- Basic method: Memory is divided into segments



- Furthermore
  - logical addresses are mapped to physical addresses
    - provides sharing, etc.
  - hardware support for address mapping
  - a memory protection violation is detected
    - user process *traps* to (interrupts) the OS

9/11/2007

9

## Protection Mechanisms (3): Timers

---

- OS code can enforce policies only if it gets a chance to run
- Timers maintain a count of elapsed (system) clock ticks
  - when timer expires, the CPU is interrupted → run the OS code
- Used for
  - interrupting hung processes
  - context switching in time-shared systems
- Access to timers is (usually) privileged
  - WHY?

9/11/2007

10

## Outline

---

### Announcements

- TA: Jay Chen (jchen@cs), 7<sup>th</sup> Floor 715 B'way
  - Office hours: Mondays 5:00 – 6:00pm, Thursdays 4:00 – 5:00pm
  - Primary point of contact for questions about Nachos labs
- Computer system structures
  - (Review) I/O, Storage, Protection
- Operating system structures
- Processes
  - process scheduling
- Threads
  - Multithreading models
- Process cooperation
  - Shared memory vs. message passing

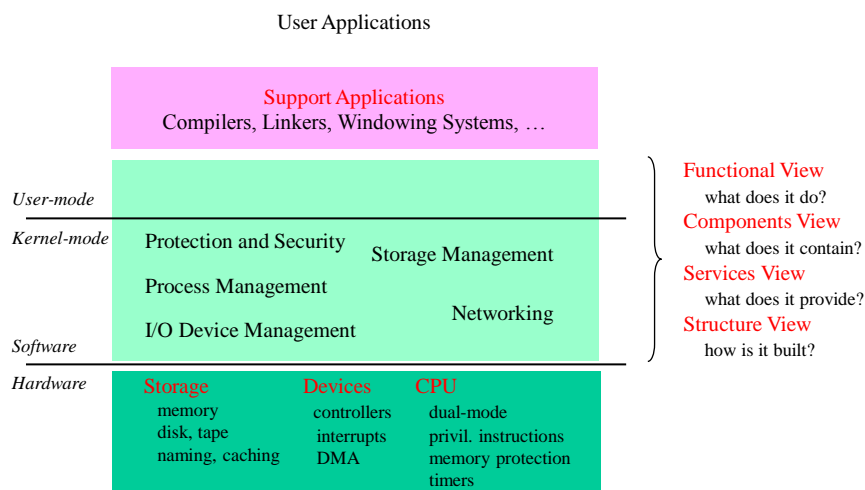
[ Silberschatz/Galvin/Gagne: Chapters 1, 2, 3.1-3.3, 4.1-4.5, 6.1 ]

9/11/2007

11

## Hardware and OS Structures

---



9/11/2007

12

## OS Views (1): Functional View

---

- What are the functions performed by an OS?
- Explicit operations
  - program execution and handling
  - I/O operations
  - file-system management
  - inter-process communication
  - exception detection and handling
    - e.g., notifying user that printer is out of paper
- Implicit operations
  - resource allocation
  - accounting
  - protection
    - e.g., maintaining data integrity, logging invalid login attempts

9/11/2007

13

## OS Views (2): Components View

---

- |  |  |
|--|--|
| <ul style="list-style-type: none"> <li>• Processes: run-time representations of user programs               <ul style="list-style-type: none"> <li>– create, terminate, suspend, resume</li> <li>– access to shared resources (e.g., printers)</li> </ul> </li> <li>• Storage               <ul style="list-style-type: none"> <li>– allocation of memory among resident processes</li> <li>– disk management (e.g., scheduling of disk accesses)</li> </ul> </li> <li>• I/O               <ul style="list-style-type: none"> <li>– device drivers, handling of device interrupts</li> <li>– files and directories</li> </ul> </li> <li>• Protection               <ul style="list-style-type: none"> <li>– user access to system resources</li> </ul> </li> </ul> | <div style="display: flex; align-items: center;"> <div style="font-size: 4em; margin-right: 10px;">}</div> <div> <p>Lectures 2-9</p> <p>Lectures 10-13</p> <p>Lectures 14-15</p> </div> </div> |
|--|--|

► Course organization follows this view

9/11/2007

14

## OS Views (3): Services View

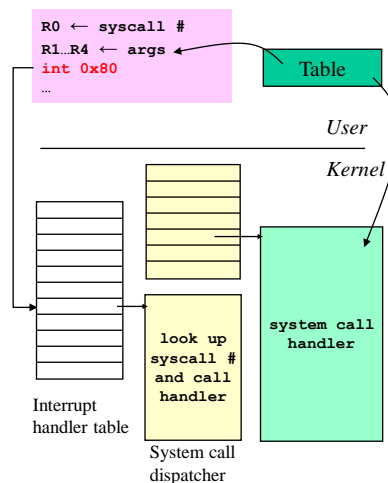
- Two issues
    - What services does an OS provide? (same as functional view)
    - How do users and user programs access these services?
  - Interface between the **user** and the OS: **Command Interpreter**
    - typical commands
      - process creation and (implicitly) destruction
      - I/O handling and file system manipulation
      - communication: interact with remote devices
      - protection management: changing file/directory access control, etc.
    - different varieties
      - the interpreter contains the code for the requested command (e.g., **delete**)
      - the interpreter calls a system routine to handle the request
      - the interpreter spawns new process(es) to handle the request
        - process lookup through some general procedure
- you will implement a simple shell in Nachos Lab 5

9/11/2007

15

## OS Views (3): Services View (contd.)

- Interface between a **user program** and the OS: **System Calls**
  - arguments passed in registers, a memory block, or on the stack
  - entry into the kernel using the *trap* mechanism
- Standard system calls
  - process control
  - file manipulation
  - device manipulation
  - information maintenance
    - *get/set* system data (time, memory/cpu usage), process and device attributes
  - communications



9/11/2007

16



## OS Views (4): Structure View

---

- How to structure OS functionality
    - Layering
    - Microkernels
    - Virtual machines
  - Designing and implementing an OS
- 
- Read Sections 2.6-2.9, Silberschatz, Galvin, and Gagne
  - Look at Nachos source code
    - Thomas Narten's roadmap

9/11/2007

17

## Outline

---

### Announcements

- TA: Jay Chen (jchen@cs), 7<sup>th</sup> Floor 715 B'way
  - Office hours: Mondays 5:00 – 6:00pm, Thursdays 4:00 – 5:00pm
  - Primary point of contact for questions about Nachos labs
- Computer system structures
  - (Review) I/O, Storage, Protection
- Operating system structures
- Processes
  - process scheduling
- Threads
  - Multithreading models
- Process cooperation
  - Shared memory vs. message passing


[ Silberschatz/Galvin/Gagne: Chapters 1, 2, 3.1-3.3, 4.1-4.5, 6.1 ]

9/11/2007

18

## What is a Process?

- A process is a **program in execution**.
- The components of a process are:
  - the program to be executed
  - the data on which the program will execute
  - the resources required by the program—such as memory and file(s)
  - the status of the execution
- A process is the unit of
  - resource ownership
  - protection
  - dispatching

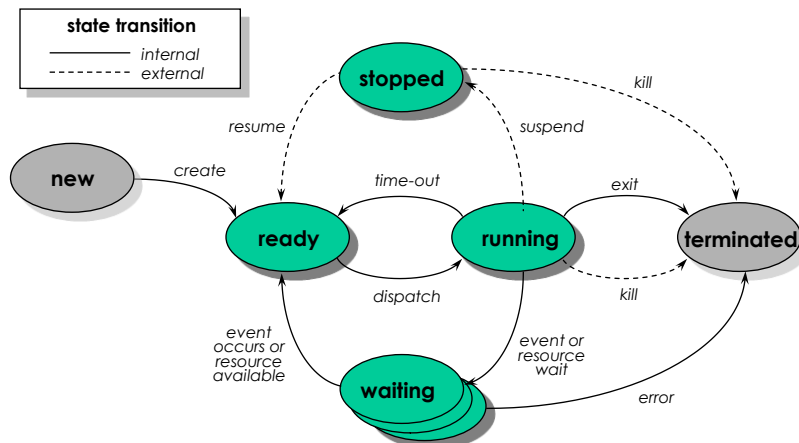


9/11/2007

19

## The State of a Process

- Can be one of: **New, Ready, Running, Waiting, Stopped, Terminated**



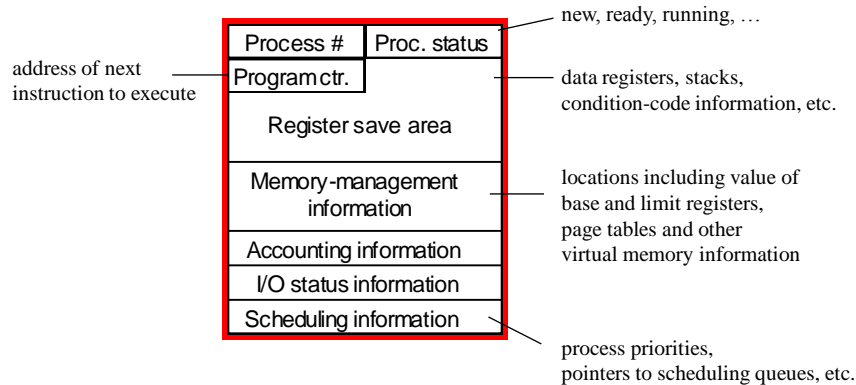
9/11/2007

20

## Process Control Information

---

### Process Control Block (PCB)



9/11/2007

21

## Scheduling Processes

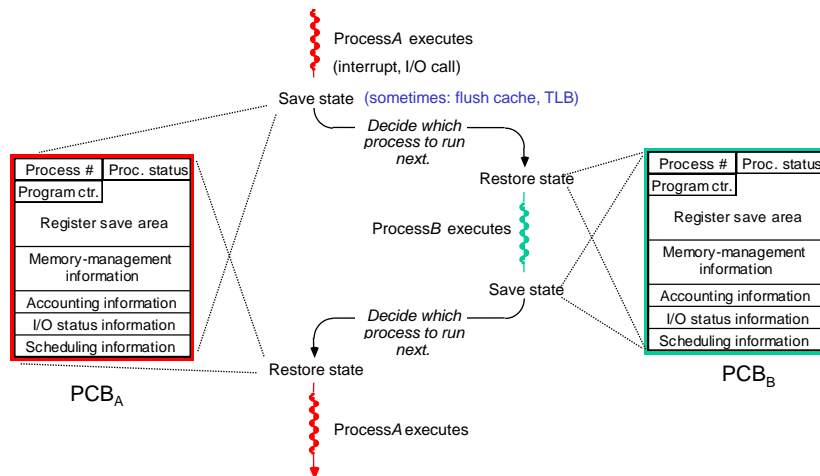
---

- “Decide which process to run next”
- Some reasons for doing this
  - move a *running* process to a *waiting* state in a multiprogrammed OS
    - multiplex CPU among ready processes
  - swapping in a time-shared system when a process' *time-slice* is over
    - typically controlled by a *timer* (process)
  - start and stop processes for accessing secondary memory and I/O
    - this may cause *spawning* of appropriate new processes
- Three main concerns
  - what happens to the process currently using the CPU?
  - how do you keep track of what each process should be doing?
  - how do you decide which process does what?

9/11/2007

22

## Concern 1: Process Context Switch

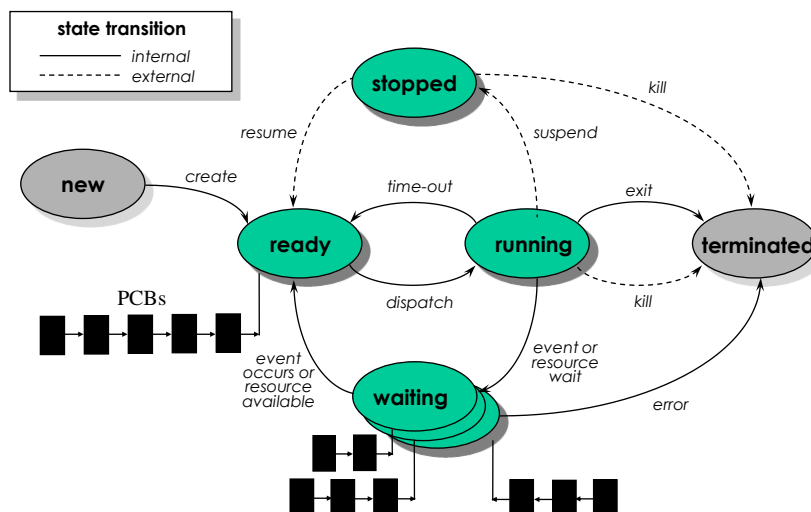


Look at the Nachos code: Thread::Yield, SWITCH  
 Nachos Lab 1: Consequences of asynchronous context switches

9/11/2007

23

## Concern 2: Process Queues



9/11/2007

24

## Concern 3: Schedulers

---

- The **long-term** scheduler
  - *operation*: creates processes and adds them to the ready queue
  - *frequency*: infrequent, ~minutes
  - *objective*: maintain good throughput by ensuring mix of I/O and CPU jobs
- The **short-term** scheduler
  - *operation*: allocates CPU and other resources to ready jobs
  - *frequency*: frequent, ~100 ms (a context switch takes ~10s of  $\mu$ secs)
  - *objective*: ensure good response times in time-sharing systems
- The **medium-term** scheduler
  - *operation*: swaps some processes out of the short-term scheduler's loop
  - *frequency*: somewhere between the short- and long-term schedulers
  - *objective*: to prevent over-multiprogramming (thrashing)
    - required when the long-term scheduler underestimates process requirements

9/11/2007

25

## System Calls for Process Management

---

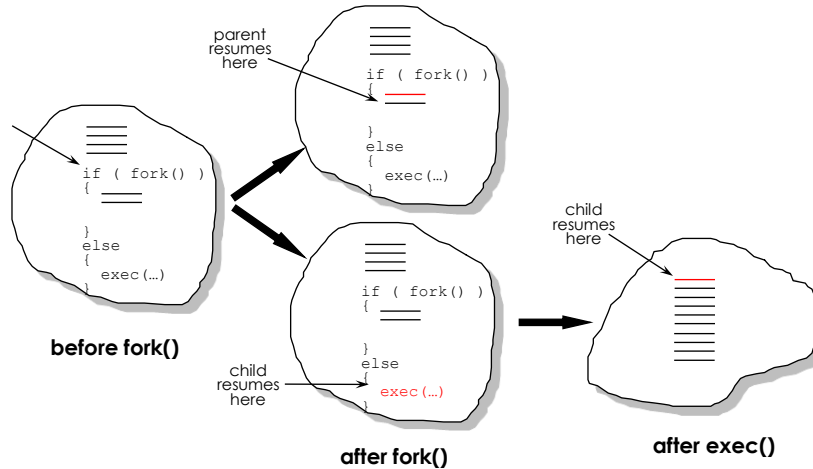
- **Creation**
  - a “parent” process spawns a “child” process; a *fork* in UNIX
    - child may or may not inherit parent's memory
    - child is added to the ready queue
  - the parent-child association is maintained via process IDs (PIDs)
- **Termination**
  - normal: a process asks the OS to delete it; an *exit* in UNIX
    - all resources of a terminated process are deallocated and reclaimed
    - on termination, the child's PID and output may be passed back to the parent
  - abnormal: another process (typically the parent) can cause termination
    - if the child exceeds its usage, becomes obsolete, or the parent is exiting the system due to some other problem
    - a process (almost always) terminates when its parent does
- **Communication**: Lecture 5
- **Coordination**: Lectures 6, 9-11

9/11/2007

26

## Example: Process Creation in UNIX

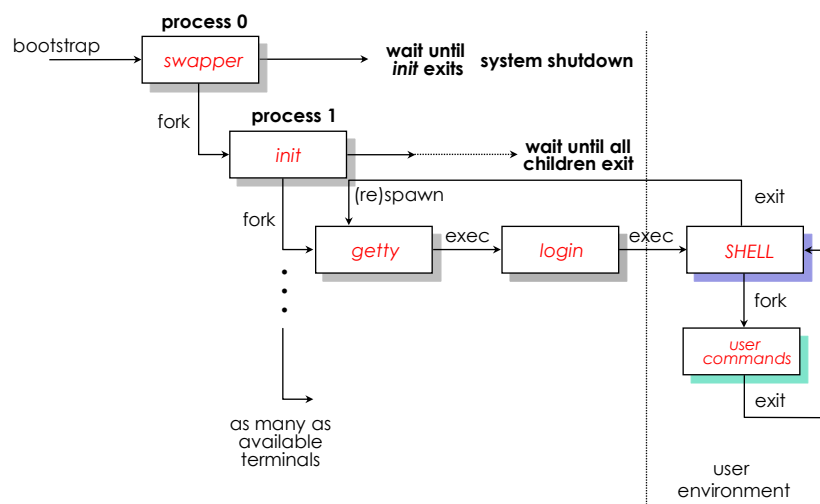
Two system calls: **fork**, **exec**



9/11/2007

27

## UNIX System Initialization



9/11/2007

28

## Outline

---

### Announcements

- TA: Jay Chen (jchen@cs), 7<sup>th</sup> Floor 715 B'way
  - Office hours: Mondays 5:00 – 6:00pm, Thursdays 4:00 – 5:00pm
  - Primary point of contact for questions about Nachos labs
- Computer system structures
  - (Review) I/O, Storage, Protection
- Operating system structures
- Processes
  - process scheduling
- **Threads**
  - Multithreading models
- **Process cooperation**
  - Shared memory vs. message passing

[ Silberschatz/Galvin/Gagne: Chapters 1, 2, 3.1-3.3, 4.1-4.5, 6.1]

9/11/2007

29

## Threads

---

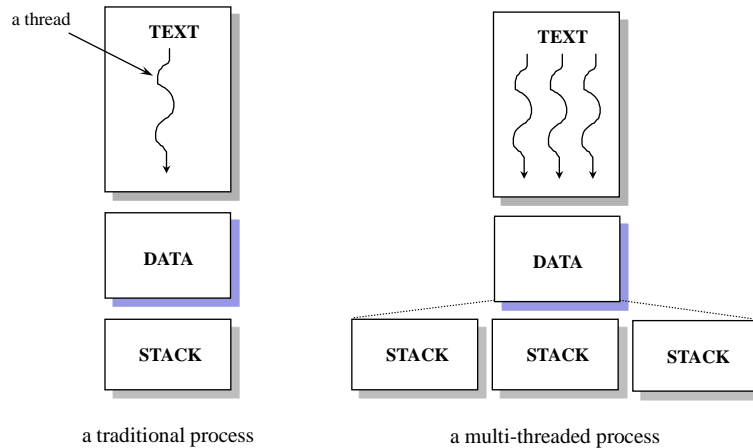
- A thread is similar to a process
  - sometimes called a *lightweight process*
  - several threads (of control) can execute within the same address space
- Like a process, a thread
  - is a basic unit of CPU utilization
  - represents the state of a program
  - can be in one of several states: *ready*, *blocked*, *running*, or *terminated*
  - has its own program counter, registers, and stack
  - executes sequentially, can create other threads, block for a system call
- Unlike a process, a thread
  - shares with peer threads, its code section, data section, and operating-system resources such as open files and signals
  - is *simpler and faster*

9/11/2007

30

## Threads versus Processes (contd.)

---



9/11/2007

31

## Threads: Why Simpler?

---

Threads share the process address space

- Benefits for the user:
  - communication is easier
  - communication is more efficient
  - security may not be necessary
    - assumed to operate within the same protection domain
  - one blocking thread need not block other threads in the process
- Benefits for the OS
  - context switching is more efficient
    - memory mappings can remain unchanged
    - cache need not be flushed
  - can run a process across multiple nodes of a multiprocessor
    - performance advantages if threads can execute in parallel (e.g., web servers)

9/11/2007

32



## Types of Threads

---

- User-level threads (e.g., **pthread**s: Section 4.3.1, Java threads: Section 4.3.3)
  - OS does not know about them
  - implemented/scheduled by library routines
  - ⬆ operations are faster (context switch, communication, control)
  - ⬇ blocking operations block the entire process (even with ready threads)
  - ⬇ operations based on local criteria may be less effective (e.g., scheduling)
- Kernel-level threads (e.g., **Solaris 2**, WinXP: Section 4.5.1, Linux: 4.5.2)
  - known to the OS
  - scheduled by the OS
  - ⬆ process need not block if one of its threads blocks on a system call
  - ⬇ thread operations are expensive
    - switching threads involves kernel interaction (via an interrupt)
  - ⬆ the kernel can do a better job of allocating resources

9/11/2007

33




## Multithreading Models

---

- Most systems provide support for both user and kernel threads

Three dominant models for mapping threads to kernel resources

- **Many-to-one**
  - Thread management done in user space
  - Entire process blocks if a thread does a blocking operation
  - E.g., systems without kernel threads
- **One-to-one**
  - Each user-thread mapped to a kernel thread
  - Allows more concurrency
  - E.g., Windows 2000, XP (**fibers**: many-to-one)
- **Many-to-many**
  - Combination of the above two
  - E.g., Solaris 2



9/11/2007

34


## POSIX Threads (pthreads)

- A portable API for multithreaded programs
  - Some pthreads implementations do map threads to kernel threads
  - Most rely on user-level threading support
    - Assembly instructions to save/restore registers
- Calls for creating, exiting, joining pthreads
  - **pthread\_create**: start execution of this thread
    - Takes function pointer as an argument
  - **pthread\_exit**: terminate execution of this thread
  - **pthread\_join**: wait for a particular thread to exit
- Other calls
  - Help set thread attributes (stack size, scheduling behavior, etc.)
  - Specify **signal handling**
    - Signals are a way of allowing processes to respond to events
      - Interrupts (Ctrl-C), others
    - Multithreaded systems need to define a way for signals to be communicated to individual threads (see Section 4.4.3)
      - All threads, a specific thread, only those threads that do not block the signal, ...

9/11/2007

35

## Processes and Threads in Solaris 2



- OS schedules execution of **kernel threads** (KTs)
    - runs them on the CPUs
    - a KT can be pinned to a CPU
  - A task consists of one or more **lightweight processes** (LWPs)
    - LWPs in a task may
      - contain several *user-level threads*
      - issue a system call
      - block
  - A LWP is associated with a KT
  - There are KTs with no LWP
- 
- Linux has a simpler model (see Section 4.5.2): only kernel threads (tasks)
    - System calls for process creation (fork) and thread creation (clone)

9/11/2007

36

## Outline

---

### Announcements

- TA: Jay Chen (jchen@cs), 7<sup>th</sup> Floor 715 B'way
  - Office hours: Mondays 5:00 – 6:00pm, Thursdays 4:00 – 5:00pm
  - Primary point of contact for questions about Nachos labs
- Computer system structures
  - (Review) I/O, Storage, Protection
- Operating system structures
- Processes
  - process scheduling
- Threads
  - Multithreading models
- Process cooperation
  - Shared memory vs. message passing

[ Silberschatz/Galvin/Gagne: Chapters 1, 2, 3.1-3.3, 4.1-4.5, 6.1]

9/11/2007

37

## Process Cooperation

---

- Why do processes cooperate?
  - *modularity*: breaking up a system into several sub-systems
    - e.g.: an interrupt handler and device driver that need to communicate
  - *convenience*: users might want to have several processes share data
  - *speedup*: a single program is run as several sub-programs
- How do processes cooperate?
  - communication abstraction: *producers* and *consumers*
    - *producers* produce a piece of information
    - *consumers* use this information
  - abstraction helps deal with general “phenomena” and simplifies correctness arguments
- Two general classes of process cooperation techniques
  - shared memory
  - message passing

9/11/2007

38

## Shared Memory (Procedure-oriented System)

- Processes can directly access data written by other processes
  - examples: POSIX threads, Java, Mesa, small multiprocessors
- A finite-capacity shared buffer

```

N: integer                                -- buffer size
nextin = nextout = 1 initially;           -- start of buffer
buffer: array of size N

Producer:
Repeat
  -- produce an item in tempin
  while (nextin+1) mod n = nextout do wait-a-bit;
  buffer[nextin] := tempin;
  nextin := (nextin+1) mod n;

Consumer:
Repeat
  while nextin = nextout do wait-a-bit;
  tempout := buffer[nextout];
  nextout := (nextout+1) mod n;
  -- consume the item in tempout

```


9/11/2007

39

## Message Passing (Message-oriented System)

- Execution is in separate address spaces
  - communication using message channels
  - examples: UNIX processes, large multiprocessors, etc.

- Components
  - **messages** and message identifiers
  - message **channels** and **ports**
    - channels (**pipes**) must be bound to ports
    - queues associated with ports
  - message transmission operations
    - SendMessage[channel, body] returns id
    - AwaitReply[id]
    - RecvMessage[port] returns id
    - SendReply[id, body]



- Many variants: See Section 4.5
- ➡ Focus on shared memory for next few lectures

9/11/2007

40

## Bounded Buffers Using Counters

```

N: integer                                -- buffer size
counter: integer = 0 initially;
nextin = nextout = 1 initially;  -- start of buffer
buffer: array of size N

Producer:
Repeat
  -- produce an item in tempin
  while counter = N do wait-a-bit;
  buffer[nextin] := tempin;
  nextin := (nextin+1) mod n;
  counter := counter+1;
Consumer:
Repeat
  while counter = 0 do wait-a-bit;
  tempout := buffer[nextout];
  nextout := (nextout+1) mod n;
  counter := counter-1;
  -- consume the item in tempout

```

Producer and Consumer processes are asynchronous! execution of these two statements can be interleaved (e.g., because of interrupts)

9/11/2007

41

## Interleaving of Increment/Decrement

- Each of increment and decrement are actually implemented as a series of machine instructions on the underlying processor

Producer	Consumer
register1 := counter	register2 := counter
register1 := register1 + 1	register2 := register2 - 1
counter := register1	counter := register2

- An interleaving
  - counter = 5; a producer followed by a consumer

Producer	Consumer	
register1 := counter		{register1 = 5}
register1 := register1 + 1		{register1 = 6}
	register2 := counter	{register2 = 5}
	register2 := register2 - 1	{register2 = 4}
counter := register1		{counter = 6}
	counter := register2	{counter = 4}

9/11/2007

42

## The Problem

---

- Increment and decrement are not *atomic* or *uninterruptable*
  - two or more operations are executed **atomically** if the result of their execution is equivalent to that of some serial order of execution
  - operations which are always executed atomically are called **atomic**
    - byte read; byte write;
    - word read; word write
- The code containing these operations creates a **race condition**
  - produces inconsistencies in shared data
- Reasons for non-atomic execution
  - interrupts
  - context-switches

9/11/2007

43

## The Solution

---

- The producer and consumer processes need to **synchronize**
  - so that they do *not* access shared variables at the same time
  - this is called **mutual exclusion**
    - the *shared* and *critical* variables can be accessed by only one process at a time
  - access must be **serialized** even if the processes attempt **concurrent** access
    - in the previous example: counter increment and decrement operations
- General framework for achieving this: **Critical Sections**
  - work independent of the particular context or need for synchronization

9/11/2007

44

## Critical Sections

---

- Critical sections: General framework for process synchronization

**ENTRY-SECTION**

**CRITICAL-SECTION-CODE**

**EXIT-SECTION**

- the **ENTRY-SECTION** controls access to make sure that no more than one process  $P_i$  gets to access the critical section at any given time
  - acts as a *guard*
- the **EXIT-SECTION** does bookkeeping to make sure that other processes that are waiting know that  $P_i$  has exited
- How can we implement critical sections?
  - turn off interrupts around critical operations
  - ✓ build on top of atomic memory load/store operations
  - ✓ provide higher-level primitives

9/11/2007

45