

G22.2250-001

Operating Systems

Lecture 4

Process Synchronization (cont'd) Classical Synchronization Problems

September 25, 2007

Outline

Announcements

- Lab 1 was due on Friday
 - Please see me after class if you have not handed this in
- Process synchronization primitives
 - (Review) Locks, Semaphores
 - Condition variables
 - Implementation techniques
- Classical synchronization problems
 - Mutual exclusion, sequencing, bounded buffer
 - Readers-writers, dining philosophers
 - A larger example
- Language support for synchronization
 - Conditional critical regions
 - Monitors

[Silberschatz/Galvin/Gagne: Chapters 6.4-6.8]

(Review)

Synchronization Primitives (1): Locks (Mutexes)

- **Locks**
 - a single boolean variable **L**
 - in one of two states: **AVAILABLE**, **BUSY**
 - accessed via two *atomic* operations
 - **LOCK** (also known as **Acquire**)


```
while ( L != AVAILABLE ) wait-a-bit
L = BUSY;
```
 - **UNLOCK** (also known as **Release**)


```
L = AVAILABLE;
wake up a waiting process (if any)
```
 - process(es) waiting on a LOCK cannot “lock-out” process doing UNLOCK
- Critical sections using locks


```
LOCK( L )
CRITICAL SECTION
UNLOCK( L )
```

 - Mutual exclusion? Progress? Bounded waiting?

9/25/2007

3

(Review)

Synchronization Primitives (2): Semaphores

- **Semaphores**
 - a single integer variable **S**
 - accessed via two *atomic* operations
 - **WAIT** (sometimes denoted by **P**)


```
while S <= 0 do wait-a-bit;
S := S-1;
```
 - **SIGNAL** (sometimes denoted by **V**)


```
S := S+1;
wake up a waiting process (if any)
```
 - WAITing process(es) cannot “lock out” a SIGNALing process
- **Binary semaphores**
 - **S** is restricted to take on only the values 0 and 1
 - **WAIT** and **SIGNAL** become similar to **LOCK** and **UNLOCK**
 - are *universal* in that counting semaphores can be built out of them

9/25/2007

4

Universality of Binary Semaphores

- Implement operations on a (counting) semaphore **CountSem**
 - use binary semaphores $S1 = 1, S2 = 0$
 - integer C = initial value of counting semaphore

P (CountSem)	V (CountSem)
P (S1) ;	P (S1) ;
$C := C - 1;$	$C := C + 1;$
if ($C < 0$) then	if ($C \leq 0$) then V(S2) ;
begin V(S1) ; P(S2) ; end	else V(S1) ;
V(S1) ;	

- $S1$ ensures mutual exclusion for accessing C
- $S2$ is used to block processes when $C < 0$
- is a race condition possible after **V(S1)** but before **P(S2)**?

9/25/2007

5

Synchronization Primitives (3): Condition Variables

- Condition variables
 - an *implicit* process queue
 - three operations that *must be performed within a critical section*
 - WAIT**

```
associate self with the implicit queue
suspend self
```
 - SIGNAL**

```
wake up exactly one suspended process on queue
```

 - has no effect if there are no suspended processes
 - BROADCAST**

```
wake up all suspended processes on queue
```
- Two types based on what happens to the process doing the **SIGNAL**
 - Mesa style (Nachos uses Mesa-style condition variables)
 - SIGNAL**-ing process continues in the critical section
 - resumed process must re-enter (so, is not guaranteed to be the next one)
 - Hoare style
 - SIGNAL**-ing process immediately exits the critical section
 - resumed process now occupies the critical section

9/25/2007

6

Uses of Condition Variables

- Can be used for constructing
 - critical sections, sequencing, ...
- Primary use is for waiting on an event to happen
 - after checking that it has not already happened
 - WHY IS THIS IMPORTANT?
- Example: Three processes that need to cycle among themselves
 - `<print 0>; <print 1>; <print 2>; <print 0>; <print 1>; ...`
 - One variable: **turn**; three condition variables: **cv₀**, **cv₁**, **cv₂**
 - Process P_i executes (in a critical section)

```

while ( turn != i) WAIT(cvi)
<do the operation>
turn := (turn + 1) mod 3; SIGNAL(cvturn)

```

9/25/2007

7

Implementing the Synchronization Primitives

- Need support for **atomic** operations from the underlying hardware
 - applicable only to a small number of instructions
 - else, can implement critical sections this way

Three choices

- Use **n-process mutual-exclusion** solutions
 - complicated
- ✓ Selectively **disable interrupts** on uniprocessors
 - so, no unanticipated context switches ➡ atomic execution
 - solution adopted in Nachos (see Lab 2 for details)
- ✓ Rely on special **hardware synchronization instructions**
- Can implement one primitive in terms of another
 - Nachos Lab 2

9/25/2007

8

Implementation Choices (1): Interrupt Disabling

- Semaphores

```

P(S)
  DISABLE-INTERRUPTS
  while S <= 0 do wait-a-bit
    [ENABLE-INTERRUPTS; YIELD CPU; DISABLE-INTERRUPTS]
  S := S-1;
  ENABLE-INTERRUPTS

V(S)
  DISABLE-INTERRUPTS
  S := S+1;
  ENABLE-INTERRUPTS

```

- Drawback

- a process spins on this loop (**busy waiting**) till it can enter critical section
- can waste *substantial* amount of CPU cycles idling
 - Even if *wait-a-bit* is implemented as
 - give up CPU (i.e. put at the end of ready queue)
 - since there are still context switches
- not a very useful utilization of valuable cycles

9/25/2007

9

Efficient Semaphores

- Implement P and V differently

- maintain an explicit *wait queue* organized as a scheduler structure

```

type semaphore = record
  value: integer;
  L: list of processes;
end;

P(S):  S.value := S.value - 1;      V(S):  S.value := S.value + 1;
if ( S.value < 0 )                  if ( S.value <= 0 )
  then begin                        then begin
    add process to S.L              remove P from S.L
    block;                          wakeup(P);
  end;                              end;

```

- still need atomicity: can use previously discussed solutions
 - can have spinning but only for a small period of time (~10 instructions)
- queue enqueue/dequeue must be fair
 - not required by semantics of semaphores

9/25/2007

10

Implementation Choices (2): Hardware Support

- Rationale: Hardware instructions enable **simpler/efficient** solutions to common synchronization problems
 - disabling interrupts is a brute-force approach
 - does not work on multiprocessors
 - simultaneous disabling of all interrupts is not feasible
- Two common primitives
 - test-and-set
 - swap

9/25/2007

11

Semantics of Hardware Primitives

- **Test-and-set**
 - given boolean variables X, Y, atomically set X := Y; Y := true

```
boolean Test-and-set( boolean &target ) {
    boolean rv = target;
    target = true;
    return rv;
}
```

- **Swap**
 - atomically exchange the values of given variables X and Y
 - can emulate test-and-set

```
boolean Test-and-set( boolean &target ) {
    boolean t := true;
    swap (target, t);
    return t;
}
```

9/25/2007

12

Implementing Locks Using Test-and-Set

```
LOCK:      L : boolean := false
           while Test-and-set(lock) wait-a-bit
UNLOCK    lock := false
```

- Properties of this implementation
 - Mutual exclusion?
 - first process P_i entering critical section sets `lock := true`
 - test-and-set (from other processes) evaluates to true after this
 - when P_i exits, lock is set to false, so the next process P_j to execute the instruction will find test-and-set = false and will enter the critical section
 - Progress?
 - trivially true
 - Unbounded waiting
 - possible since depending on the timing of evaluating the test-and-set primitive, other processes can enter the critical section first
 - See Section 6.4 for a solution to this problem

9/25/2007

13

Synchronization Primitives in Real OSe

- Unix: Single CPU OS
 - implement critical sections using interrupt elevation
 - disallow interrupts that can modify the same data
 - (Linux 2.4 and earlier, Section 6.8.3) disable kernel preemption
 - another possibility: interrupts never “force” a context switch
 - they just set flags, or wake up processes
 - primitives
 - `sleep` (address);
 - `wake_up` (address); -- wakes up all processes sleeping on address
 - typical code


```
ENTRY: while (locked) sleep(bufaddr);
        locked = true;
EXIT:  locked = false; wake_up (bufaddr);
```

9/25/2007

14

Synchronization Primitives in Real OSes (contd.)

- Solaris 2: multi-CPU OS
 - for brief accesses only
 - **adaptive** mutexes
 - starts off as a standard spinlock semaphore
 - if lock is held by running thread, continues to spin
 - » valid only on a multi-CPU system
 - otherwise blocks
 - for long-held locks
 - (process queue) semaphores
 - condition variables
 - wait and signal
 - **reader-writer locks**
 - for frequent mostly read-only accesses
 - turnstiles
 - the queue structure on which threads block when waiting for a lock
 - associated with threads rather than lock objects
 - Each thread can block on at most one object, so more efficient

9/25/2007

15

Outline

Announcements

- Lab 1 was due on Friday
 - Please see me after class if you have not handed this in
- Process synchronization primitives
 - (Review) Locks, Semaphores
 - Condition variables
 - Implementation techniques
- Classical synchronization problems
 - Mutual exclusion, sequencing, bounded buffer
 - Readers-writers, dining philosophers
 - A larger example
- Language support for synchronization
 - Conditional critical regions
 - Monitors

[Silberschatz/Galvin/Gagne: Chapters 6.4-6.8]

9/25/2007

16

Classical Synchronization Problems

- Commonly encountered problems in operating systems
 - used to test any proposal for a new synchronization primitive
- 1. **Mutual exclusion**
 - only one process executes a piece of code (critical section) at any time
 - OS examples: access to shared resources
 - e.g., a printer
- 2. **Sequencing**
 - a process waits for another process to finish executing some code
 - OS examples: waiting for an event
 - e.g., `recv` suspends until there is some data to read on the network

9/25/2007

17

Classical Synchronization Problems (cont'd)

- 3. **Bounded-buffer** (also referred to as the **Producer-Consumer** problem)
 - a pool of n buffers
 - *producer* process(es) put items into the pool
 - *consumer* process(es) take items out of the pool
 - issues: mutual exclusion, empty pool, and full pool
 - OS examples: buffering for pipes, file caches, etc.
- 4. **Readers-Writers**
 - multiple processes access a shared data object X
 - any number of *readers* can access X at the same time
 - no *writer* can access it at the same time as a *reader* or another *writer*
 - mutual exclusion is too constraining: WHY?
 - variations:
 - **reader-priority**: a reader must not wait for a writer
 - **writer-priority**: a writer must not wait for a reader
 - OS examples: file locks


9/25/2007

18

Classical Synchronization Problems (contd.)

5. Dining Philosophers

- 5 philosophers
- 5 chopsticks placed between them
 - to eat requires two chopsticks
- philosophers alternate between thinking and eating
- issues: **deadlock**, **starvation**, **fairness**
- OS examples: simultaneous use of multiple resources
 - e.g., disk bandwidth and storage



9/25/2007

19

Mutual Exclusion and Sequencing Using Semaphores

- Mutual exclusion: Semaphore initialized to 1

```

P(S) ;
CRITICAL SECTION
V(S) ;

```

- Sequencing: Semaphore initialized to 0

<i>process 1</i>	<i>process 2</i>
	B() ;
	V(S) ;
P(S) ;	
A() ;	

9/25/2007

20

Bounded-buffer Using Semaphores

- Three semaphores
 - mutex**: provide mutual exclusion between processes (initial value = 1)
 - empty**: count the number of empty slots (initial value = N)
 - full**: count the number of full slots (initial value = 0)

Producer(s) :

```
repeat
  // produce item in nextp
  P( empty );
  P( mutex );
  // add nextp to buffer
  V( mutex );
  V( full );
until false;
```

Consumer(s) :

```
repeat
  P( full );
  P( mutex );
  // remove item to nextc
  V( mutex );
  V( empty );
  // consume item in nextc
until false;
```

9/25/2007

21

Readers-Writers Using Semaphores

To allow multiple readers, synchronize only the first/last reader with writers

Reader(s)

```
P(x);
rcount := rcount + 1;
if (rcount == 1) then P(wsem);
V(x);

READ

P(x);
rcount := rcount - 1;
if (rcount == 0) then V(wsem);
V(x);
```

Writer(s)

```
P(wsem);
WRITE
V(wsem);
```

stream of readers
can starve writers

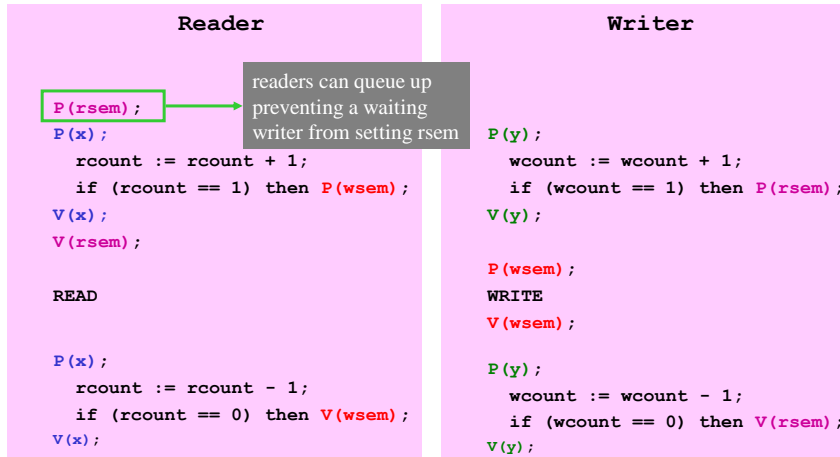
can release either
waiting readers
or writers

9/25/2007

22

Readers-Writers Using Semaphores: Writer-Priority

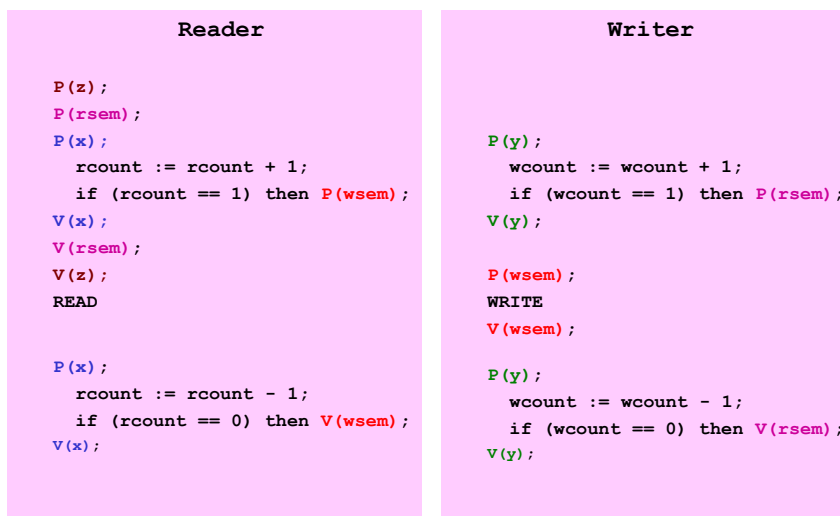
Have a writer block out subsequent readers (same as readers block out writers)



9/25/2007

23

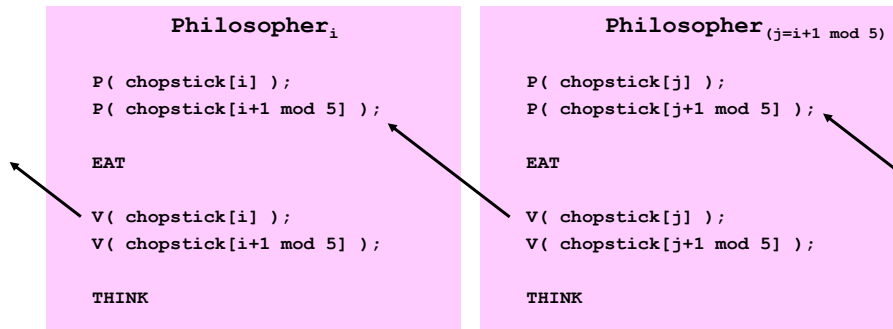
Readers-Writers Using Semaphores: Writer-Priority (2)



9/25/2007

24

Dining Philosophers Using Semaphores

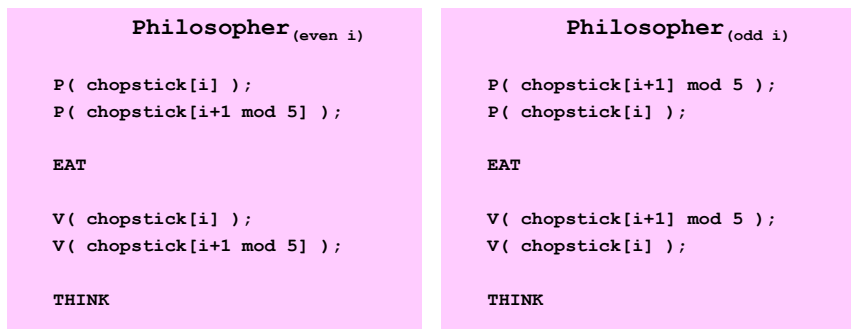


- Deadlock
 - a set of processes is in a deadlock state when every process in the set is waiting for an event that can be caused only by another process in the set*
 - details in Lectures 5 and 6.

9/25/2007

25

Dining Philosophers Using Semaphores - 2



- Alternate solutions
 - allow at most 4 philosophers to sit simultaneously at the table
 - allow a philosopher to pick up chopsticks only if both are available
- All of these solutions suffer from the possibility of **starvation!**

9/25/2007

26


A Larger Example: A Barbershop Problem

- Example taken from
 - Operating Systems: Internals and Design Principles, 3rd Edition
 - William Stallings, Prentice Hall, 1998
- The problem: Orchestrating activities in a barbershop
 - 3 chairs, 3 barbers, 1 cash register, waiting area: 4 customers on a sofa, plus additional standing room
 - Fire codes limit total number of customers to 20 at a time
 - A customer
 - Will not enter the shop if it is filled to capacity
 - Takes a seat on the sofa, or stands if sofa is filled
 - When a barber is free, the customer waiting longest on sofa is served
 - The customer standing the longest takes up seat on the sofa
 - When a customer's haircut is finished, any barber can accept payment but because of the single cash register, only one payment is accepted at a time
 - Barbers divide their time between cutting hair, accepting payment, and sleeping

9/25/2007

27

A Barbershop Problem (cont'd)



- Shop and sofa capacity
 - **max_capacity** (initial value = 20)
 - **sofa** (initial value = 4)
- Barber chair capacity
 - **barber_chair** (initial value = 3)
- Ensuring customers are in barber chair
 - **cust_ready** (initial value = 0)
 - barber waits for customer
 - **finished** (initial value = 0)
 - customer waits for haircut to finish
 - **leave_b_chair** (initial value = 0)
 - barber waits for chair to empty
- Paying and receiving
 - **payment** (initial value = 0)
 - cashier waits for customer to pay
 - **receipt** (initial value = 0)
 - customer waits for cashier to ack
- Coordinating barber functions
 - **coord** (initial value = 0)
 - wait for a barber resource to free up

9/25/2007

28

A Barbershop Problem (cont'd)

- **Shop and sofa capacity**
 - `max_capacity` (`:= 20`)
 - `sofa` (`:= 4`)
- **Barber chair capacity**
 - `barber_chair` (`:= 3`)
- **Ensuring customers are in barber chair**
 - `cust_ready` (`:= 0`)
 - `finished` (`:= 0`)
 - `leave_b_chair` (`:= 0`)
- **Paying and receiving**
 - `payment` (`:= 0`)
 - `receipt` (`:= 0`)
- **Coordinating barber functions**
 - `coord` (`:= 0`)

Customer

```
P( max_capacity );
// enter shop
P( sofa );
// sit on sofa
P( barber_chair );
// get up from sofa
V( sofa );
// sit in barber chair
V( cust_ready );
P( finished );
// leave barber chair
V( leave_b_chair );
// pay
V( payment );
P( receipt );
// exit shop
V( max_capacity );
```

Barber

```
P( cust_ready );
P( coord );
// cut hair
V( coord );
V( finished );
// wait for customer to leave
P( leave_b_chair );
// tell next customer to hop on
V( barber_chair );
```

Cashier

```
P( payment );
P( coord );
// accept payment
V( coord );
V( receipt );
```

9/25/2007

29

A Barbershop Problem (cont'd): Mutual Exclusion

- **Shop and sofa capacity**
 - `max_capacity` (`:= 20`)
 - `sofa` (`:= 4`)
- **Barber chair capacity**
 - `barber_chair` (`:= 3`)
- **Ensuring customers are in barber chair**
 - `cust_ready` (`:= 0`)
 - `finished` (`:= 0`)
 - `leave_b_chair` (`:= 0`)
- **Paying and receiving**
 - `payment` (`:= 0`)
 - `receipt` (`:= 0`)
- **Coordinating barber functions**
 - `coord` (`:= 0`)

Customer

```
P( max_capacity );
// enter shop
P( sofa );
// sit on sofa
P( barber_chair );
// get up from sofa
V( sofa );
// sit in barber chair
V( cust_ready );
P( finished );
// leave barber chair
V( leave_b_chair );
// pay
V( payment );
P( receipt );
// exit shop
V( max_capacity );
```

Barber

```
P( cust_ready );
P( coord );
// cut hair
V( coord );
V( finished );
// wait for customer to leave
P( leave_b_chair );
// tell next customer to hop on
V( barber_chair );
```

Cashier

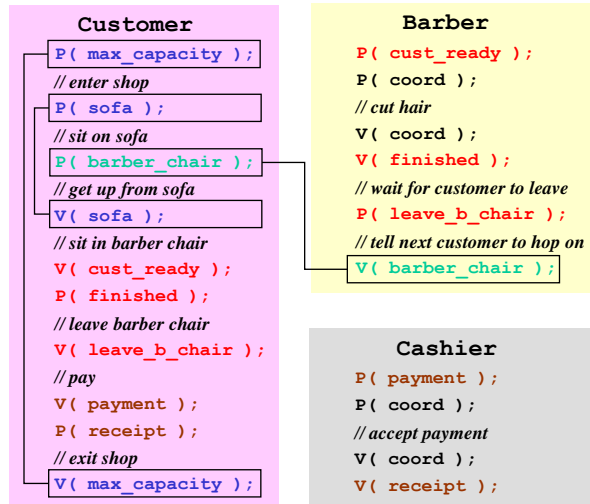
```
P( payment );
P( coord );
// accept payment
V( coord );
V( receipt );
```

9/25/2007

30

A Barbershop Problem (cont'd): Bounded Buffer

- **Shop and sofa capacity**
 - max_capacity ($:= 20$)
 - sofa ($:= 4$)
- **Barber chair capacity**
 - barber_chair ($:= 3$)
- **Ensuring customers are in barber chair**
 - cust_ready ($:= 0$)
 - finished ($:= 0$)
 - leave_b_chair ($:= 0$)
- **Paying and receiving**
 - payment ($:= 0$)
 - receipt ($:= 0$)
- **Coordinating barber functions**
 - coord ($:= 0$)

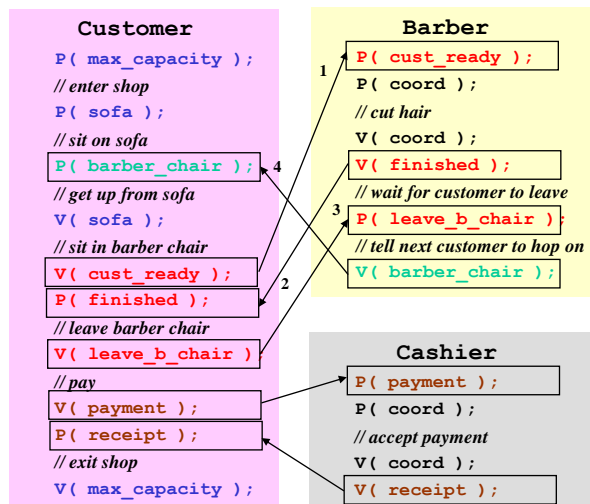


9/25/2007

31

A Barbershop Problem (cont'd): Sequencing

- **Shop and sofa capacity**
 - max_capacity ($:= 20$)
 - sofa ($:= 4$)
- **Barber chair capacity**
 - barber_chair ($:= 3$)
- **Ensuring customers are in barber chair**
 - cust_ready ($:= 0$)
 - finished ($:= 0$)
 - leave_b_chair ($:= 0$)
- **Paying and receiving**
 - payment ($:= 0$)
 - receipt ($:= 0$)
- **Coordinating barber functions**
 - coord ($:= 0$)



9/25/2007

32

A Barbershop Problem (cont'd)

- Some problems with the current solution
 - since all customers are waiting on the same semaphore (**finished**), the one who started earliest is released when a barber does **V(finished)**
 - even if the haircut is not done
 - similar problem with the cashier and the **pay** and **receipt** semaphores
 - cashier may accept money from one customer and release another
 - a customer needs to wait on the sofa even if a barber chair is free
- All of these can be solved using additional semaphores

9/25/2007

33

Outline

Announcements

- Lab 1 was due on Friday
 - Please see me after class if you have not handed this in
- Process synchronization primitives
 - (Review) Locks, Semaphores
 - Condition variables
 - Implementation techniques
- Classical synchronization problems
 - Mutual exclusion, sequencing, bounded buffer
 - Readers-writers, dining philosophers
 - A larger example
- Language support for synchronization
 - Conditional critical regions
 - Monitors

[Silberschatz/Galvin/Gagne: Chapters 6.4-6.8]

9/25/2007

34

Limitations of Semaphores

- No abstraction and modularity
 - a process that uses a semaphore has to know which other processes use the semaphore, and how these processes use the semaphore
 - a process cannot be written in isolation
- Consider sequencing between three processes
 - $P_1, P_2, P_3, P_1, P_2, P_3, \dots$

P_1	P_2	P_3
<code>P(sem₁);</code>	<code>P(sem₂);</code>	<code>P(sem₃);</code>
<code>// do stuff</code>	<code>// do stuff</code>	<code>// do stuff</code>
<code>V(sem₂);</code>	<code>V(sem₃);</code>	<code>V(sem₁);</code>

What happens if there are only two processes?

What happens if you want to use this solution for four processes?

9/25/2007

35

Limitations of Semaphores (cont'd)

- Very easy to write incorrect code
 - changing the order of P and V
 - can violate mutual exclusion requirements


```
V( mutex ); CODE; P( mutex ); instead of
P( mutex ); CODE; V( mutex );
```
 - can cause deadlock


```
P( seq ); instead of
V( seq );
```
 - similar problems with omission
- Extremely difficult to verify programs for correctness
 - Need for still higher-level synchronization abstractions!

9/25/2007

36

Language Support

- Helps simplify expression of synchronization
 - more convenient
 - more secure
 - less buggy
- We shall examine two fundamental constructs
 - conditional critical regions
 - monitors
- These constructs can be found in several concurrent languages
 - Communicating Sequential Processes (CSP) *critical regions*
 - Concurrent Pascal *monitors*
 - object-oriented languages: Modula-2, Concurrent C, Java
 - Ada83, Ada95

9/25/2007

37

Conditional Critical Regions

- A high-level language declaration
 - informally, it can be used to specify that while a statement S is being executed, no more than one process can access a distinguished variable v
 - notation

```
var v: shared t;
region v when B do S;
```

- v is shared and of type t
 - can only be accessed within a **region** statement
 - B is a Boolean expression
 - S is a statement
 - can be a compound statement
- Semantics
 - A process is guaranteed **mutually exclusive access** to the region v
 - Checking of B and entry into the region happens **atomically**

9/25/2007

38

Conditional Critical Regions: Benefits

Bounded-buffer producer/consumer

```

var buffer : shared record
  pool: array [0..n-1] of item;
  count, in, out: integer;
end;

Producer:
  region buffer when count < n
  do begin
    pool[in] := nextp;
    in := (in + 1) mod n;
    count := count + 1;
  end;

Consumer:
  region buffer when count > 0
  do begin
    nextc := pool[out];
    out := (out + 1) mod n;
    count := count - 1;
  end;

```

- Guards against simple errors associated with semaphores
 - e.g., changing the order of P and V operations, or forgetting to put one of them
- Division of responsibility
 - the *developer* does not have to program the semaphore or alternate synchronization explicitly
 - the *compiler* ``automatically" plugs in the synchronization code using predefined libraries
 - once done carefully, *reduces* likelihood of mistakes in designing the delicate synchronization code

9/25/2007

39

Conditional Critical Regions: Implementation

```

var mutex: semaphore;
var delay: semaphore;
var count: integer;

P( mutex );
while not B
do begin
  try-and-enter;
end;
S;
leave-critical-region;

count++;
V( mutex );
P( delay );
// check condition
if ( not B )
  if ( count > 1 )
    // release another
    V( delay );
  P( delay );
else
  V( mutex );
  P( delay );
else count-- ;

if ( count > 0 )
  V( delay );
else V( mutex );

var first, second: semaphore;
var fcount, scount: integer;

fcount++;
if ( scount > 0 ) V( second );
else V( mutex );
P( first );
fcount-- ;
scount++ ;
if ( fcount > 0 ) V( first );
else V( second );
P( second );
scount-- ;

if ( fcount > 0 ) V( first );
else if ( scount > 0 ) V( second );
else V( mutex );


```

9/25/2007

40

Language Support (2): Monitors

- An abstract data type
 - private data
 - public procedures
 - only one procedure can be in the monitor at one time
 - each procedure may have
 - local variables
 - formal parameters
 - condition variables
 - queues of processes
 - *wait*: block on a condition variable
 - *signal*: unblock a waiting process
 - no-op if no process is waiting
- Processes can only invoke the public procedures
 - raises the granularity of atomicity to a single user-defined procedure



9/25/2007

41

Waiting in the Monitor

- Note that the semantics of executing a *wait* in the monitor is that several processes can be waiting “inside” the monitor at any given time but only one is executing
 - wait queues are internal to the monitor
 - there can be multiple wait queues
- Who executes after a signal operation? (say P signals Q)
 - (Hoare semantics) signallee Q continues
 - logically natural since the condition that enabled Q might no longer be true when Q eventually executes
 - P needs to wait for Q to exit the monitor
 - (Mesa semantics) signaller P continues
 - Q is enabled but gets its turn only after P either leaves or executes a *wait*
 - require that the *signal* be the last statement in the procedure
 - advocated by Brinch Hansen (Concurrent Pascal)
 - easy to implement but less powerful than the other two semantics

9/25/2007

42

Use of Monitors: Bounded-buffer

```

type bounded_buffer = monitor

  var buffer: array [0..N] of char;
  var in, out, count: integer;
  var notfull, notempty: condition;

  procedure entry append ...
  procedure entry remove ...

begin
  in = 0; out = 0; count = 0;
end;

procedure entry append(x: char);
  if (count==N) notfull.wait;
  buffer[in] := x;
  in := (in+1) mod N;
  count := count+1;
  notempty.signal;

procedure entry remove(x: char);
  if (count==0) notempty.wait;
  x := buffer[out];
  out := (out+1) mod N;
  count := count-1;
  notfull.signal;

```

Is this solution correct under all monitor semantics? (P signals Q)

Hoare: Q continues, P suspends YES
 Mesa: P continues, Q is put into ready queue NO
 Brinch-Hansen: P exits monitor, Q continues YES

9/25/2007

43

Use of Monitors: Bounded-buffer (Mesa Semantics)

```

type bounded_buffer = monitor

  var buffer: array [0..N] of char;
  var in, out, count: integer;
  var notfull, notempty: condition;

  procedure entry append ...
  procedure entry remove ...

begin
  in = 0; out = 0; count = 0;
end;

procedure entry append(x: char);
  while (count==N) notfull.wait;
  buffer[in] := x;
  in := (in+1) mod N;
  count := count+1;
  notempty.signal;

procedure entry remove(x: char);
  while (count==0) notempty.wait;
  x := buffer[out];
  out := (out+1) mod N;
  count := count-1;
  notfull.signal;

```

9/25/2007

44

Use of Monitors: Dining Philosophers

- Goal: Solve DP without deadlocks
- Informally:
 - algorithm for Philosopher I


```
dp.pickup(i);
eat;
dp.putdown(i);
```
 - use array to describe state


```
var state: array [0..4] of
(thinking, hungry,
eating);
```
 - use array of condition variables to block on when required resources are unavailable


```
var self: array [0..4] of
condition;
```
- **pickup(i)**
 - changes state to hungry
 - checks if neighbors are eating
 - if not, grabs chopsticks, and changes state to eating
 - otherwise, waits on self(i)
- **putdown(i)**
 - checks both neighbors
 - if either is hungry and can proceed, releases him/her

9/25/2007

45

Dining Philosophers using Monitors - 2

```

type dining_philosophers = monitor
  var state: array [0..4] of
    (thinking, hungry, eating);
  var self: array [0..4] of
    condition;

  procedure entry pickup ...
  procedure entry putdown ...
  procedure test ...

begin
  for i := 0 to 4 do
    state[i] := thinking;
  end;

  procedure entry pickup(i: 0..4);
    state[i] := hungry;
    test(i);
    while ( state[i] != eating )
      self[i].wait;

  procedure entry putdown(i: 0..4);
    state[i] := thinking;
    test (ln(i));
    test (rn(i));

  procedure test(i: 0..4);
    if (state[ln(i)] != eating and
        state[i] == hungry and
        state(rn(i)) != eating)
      state[i] := eating;
      self[i].signal;

```

9/25/2007

46

Dining Philosophers using Monitors - 3

- What is missing?
 - philosophers cannot deadlock but can starve
 - for example, we can construct timing relationships such that a waiting philosopher will be stuck in the “self” queue forever
 - monitors have to be enhanced with a fair scheduling policy to avoid starvation
 - both at the level of accessing the monitor
 - as well as to regulate “waking-up” those that are waiting inside
 - how can this be done?
 - use fair enqueue and dequeue policies

9/25/2007

47

Monitors: Other Issues

- Expressibility: Are monitors more/less powerful than semaphores or conditional critical regions?
 - these three constructs are equivalent
 - the same kinds of synchronization problems can be expressed in each
 - the other two can be implemented using any one of the constructs
 - e.g., critical regions and monitors using semaphores
 - we talked about how critical regions can be implemented
 - in Lab 2: you built condition variables using semaphores
 - » this implementation can be extended to build monitors
- Do monitors have any limitations?
 - absence of concurrency within a monitor
 - workarounds introduce all the problems of semaphores
 - monitor procedures will need to be invoked before and after
 - possibility of improper access, deadlock, etc.

9/25/2007

48