

Storing Data: Disks and Files

Chapter 7

Unsophisticated users (Customers,
Travel agents)

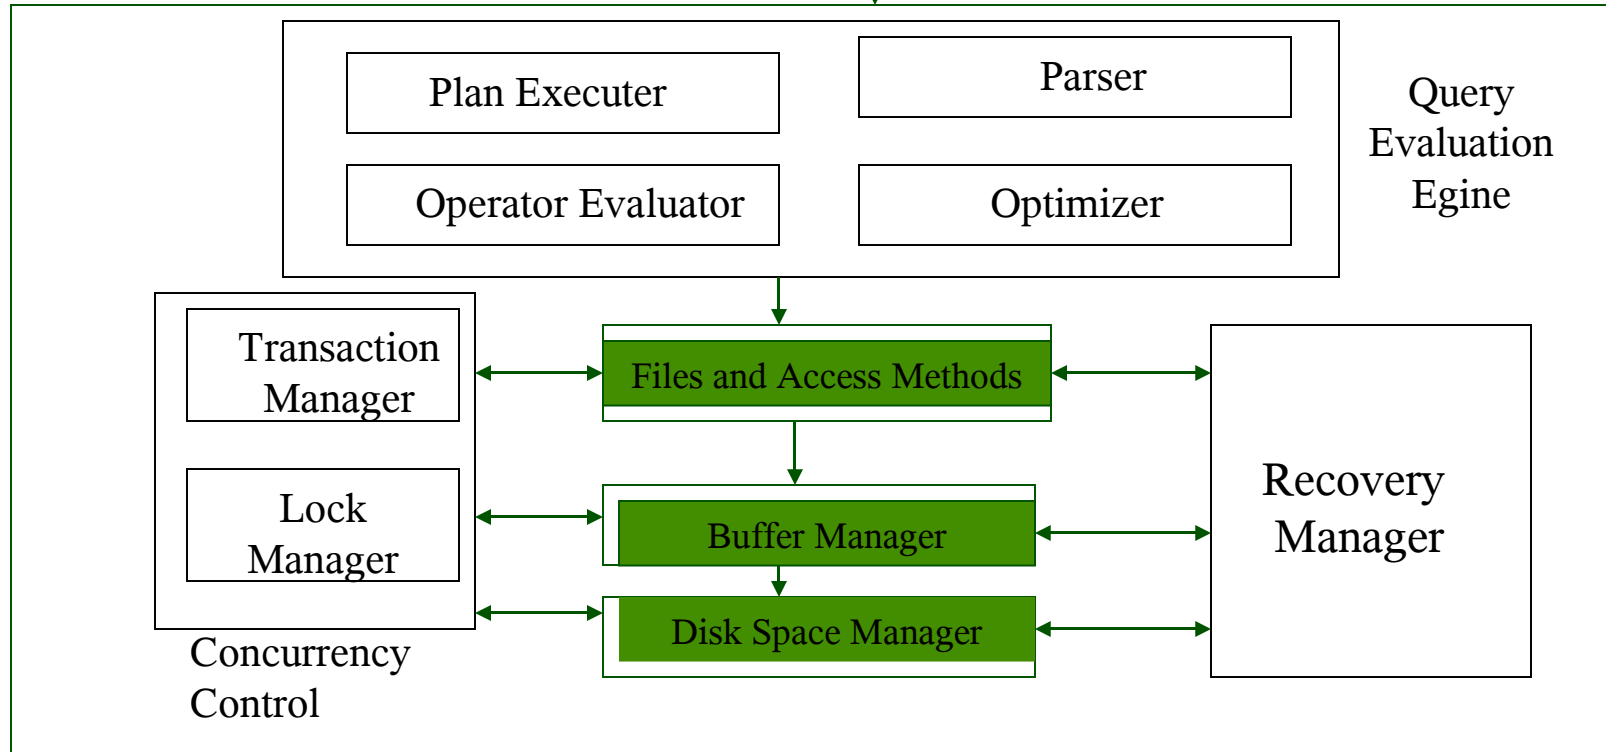
Sophisticated users, application
Programmers, DB Administrators

Web Forms


Application Front Ends

SQL Interface

SQL COMMANDS



Architecture
Of
DBMS



DBMS components covered


❖ We will cover

- disc space manager
- the buffer manager
- Implementation oriented aspects of the files and access methods layer

Outline

- ❖ The memory hierarchy
- ❖ RAID
- ❖ Disk space management
- ❖ Buffer Manager
- ❖ Files of records
- ❖ Page formats
- ❖ Record formats

The Memory Hierarchy



The Memory Hierarchy

- ❖ Primary Storage
 - Cache and main memory
 - Provides fast access to data
- ❖ Secondary storage
 - Disk
- ❖ Tertiary storage
 - Slow
 - Tapes and optical disks
- ❖ Cost of main memory is 100 times of main memory
- ❖ Tapes play a major role as data size is huge.
- ❖ We have to build Database systems to retrieve data from lower levels of memory hierarchy into main memory.

Disks and Files

- ❖ DBMS stores information on (“hard”) disks.
- ❖ This has major implications for DBMS design!
 - **READ:** transfer data from disk to main memory (RAM).
 - **WRITE:** transfer data from RAM to disk.
 - Both are high-cost operations, relative to in-memory operations, so must be planned carefully!

Why Not Store Everything in Main Memory?

- ❖ *Costs too much.* \$1000 will buy you either 128MB of RAM or 7.5GB of disk today.
- ❖ *Main memory is volatile.* We want data to be saved between runs. (Obviously!)
- ❖ Typical storage hierarchy:
 - Main memory (RAM) for currently used data.
 - Disk for the main database (secondary storage).
 - Tapes for archiving older versions of the data (tertiary storage).

Disks

- ❖ Secondary storage device of choice.
- ❖ Main advantage over tapes: random access vs. *sequential*.
- ❖ Data is stored and retrieved in units called *disk blocks* or *pages*.
- ❖ Unlike RAM, time to retrieve a disk page varies depending upon location on disk.
 - Therefore, relative placement of pages on disk has major impact on DBMS performance!


Components of a Disk

- ❖ The platters spin (say, 90rps).

- ❖ The arm assembly is moved in or out to position a head on a desired track. Tracks under heads make a *cylinder* (imaginary!).

- ❖ Only one head reads/writes at any one time.

- ❖ *Block size* is a multiple of *sector size* (which is fixed).



Accessing a Disk Page

- ❖ Time to access (read/write) a disk block:
 - *seek time* (moving arms to position disk head on track)
 - *rotational delay* (waiting for block to rotate under head)
 - *transfer time* (actually moving data to/from disk surface)
- ❖ Seek time and rotational delay dominate.
 - Seek time varies from about 1 to 20msec
 - Rotational delay varies from 0 to 10msec
 - Transfer rate is about 1msec per 4KB page
- ❖ Key to lower I/O cost: **reduce seek/rotation delays!**
Hardware vs. software solutions?

Arranging Pages on Disk

- ❖ *'Next'* block concept:
 - blocks on same track, followed by
 - blocks on same cylinder, followed by
 - blocks on adjacent cylinder
- ❖ Blocks in a file should be arranged sequentially on disk (by *'next'*), to minimize seek and rotational delay.
- ❖ For a *sequential scan*, *pre-fetching* several pages at a time is a big win!

Outline

- ❖ The memory hierarchy
- ❖ RAID
- ❖ Disk space management
- ❖ Files of records
- ❖ Page formats
- ❖ Record formats

RAID

- ❖ Disk Array: Arrangement of several disks that gives abstraction of a single, large disk.
- ❖ Goals: Increase performance and reliability.
- ❖ Two main techniques:
 - **Data striping:** Data is partitioned; size of a partition is called the striping unit. Partitions are distributed over several disks.
 - **Redundancy:** More disks \Rightarrow more failures. Redundant information allows reconstruction of data if a disk fails.

RAID Levels

- ❖ Level 0: No redundancy
- ❖ Level 1: Mirrored (two identical copies)
 - Each disk has a mirror image (check disk)
 - Parallel reads, a write involves two disks.
 - Maximum transfer rate = transfer rate of one disk
- ❖ Level 0+1: Striping and Mirroring
 - Parallel reads, a write involves two disks.
 - Maximum transfer rate = aggregate bandwidth

RAID Levels (Contd.)

❖ Level 3: Bit-Interleaved Parity

- Striping Unit: One bit. One check disk.
- Each read and write request involves all disks; disk array can process one request at a time.

❖ Level 4: Block-Interleaved Parity

- Striping Unit: One disk block. One check disk.
- Parallel reads possible for small requests, large requests can utilize full bandwidth
- Writes involve modified block and check disk

❖ Level 5: Block-Interleaved Distributed Parity

- Similar to RAID Level 4, but parity blocks are distributed over all disks

Outline

- ❖ The memory hierarchy
- ❖ RAID
- ❖ Disk space management
- ❖ Buffer Manager
- ❖ Files of records
- ❖ Page formats
- ❖ Record formats

Disk Space Management

- ❖ Lowest layer of DBMS software manages space on disk.
- ❖ Supports the concept of a page as a unit of data.
- ❖ Higher-levels issue commands
 - allocate/de-allocate a page
 - read/write a page
- ❖ Request for a *sequence* of pages must be satisfied by allocating the pages sequentially on disk! Higher levels don't need to know how this is done, or how free space is managed.
- ❖ The disk space manager hides details of the underlying hardware.
- ❖ Higher levels of the software think of the data as a collection of pages.

Keeping the track of Free Blocks

- ❖ Although blocks are allocated sequentially on disk, allocations and de-allocations could in general create 'holes'
- ❖ Methods
 - Linked-list approach: Keep the list of free blocks and occupied blocks.
 - Maintain a bit-map with one bit for each disk block.
 - Bit-map method allows very fast identification and allocation of contiguous areas on disk.

Using OS File Systems to Manage Disk Space

❖ OS

- supports the abstraction of a file as a sequence of bytes.
- Manages space on disk and translates requests, such as “Read byte *i* of file *f*, into corresponding low level instructions: “Read block *m* of track *t* of cylinder *c* of disk *d*”.

❖ DB disk manager can be build using OS files

- Entire DB can be managed with one or more OD files.
- The disk space manager is responsible for managing the space in these OS files.

❖ Many DB systems do not rely on the OS file system

- They have separate disk management
- Reason: DBMS can not be specific to any OS file system
- DBMS may want to access bigger files than the file sizes OS supports
- OS file may not span more than one disk, but in DBMS it is necessary.

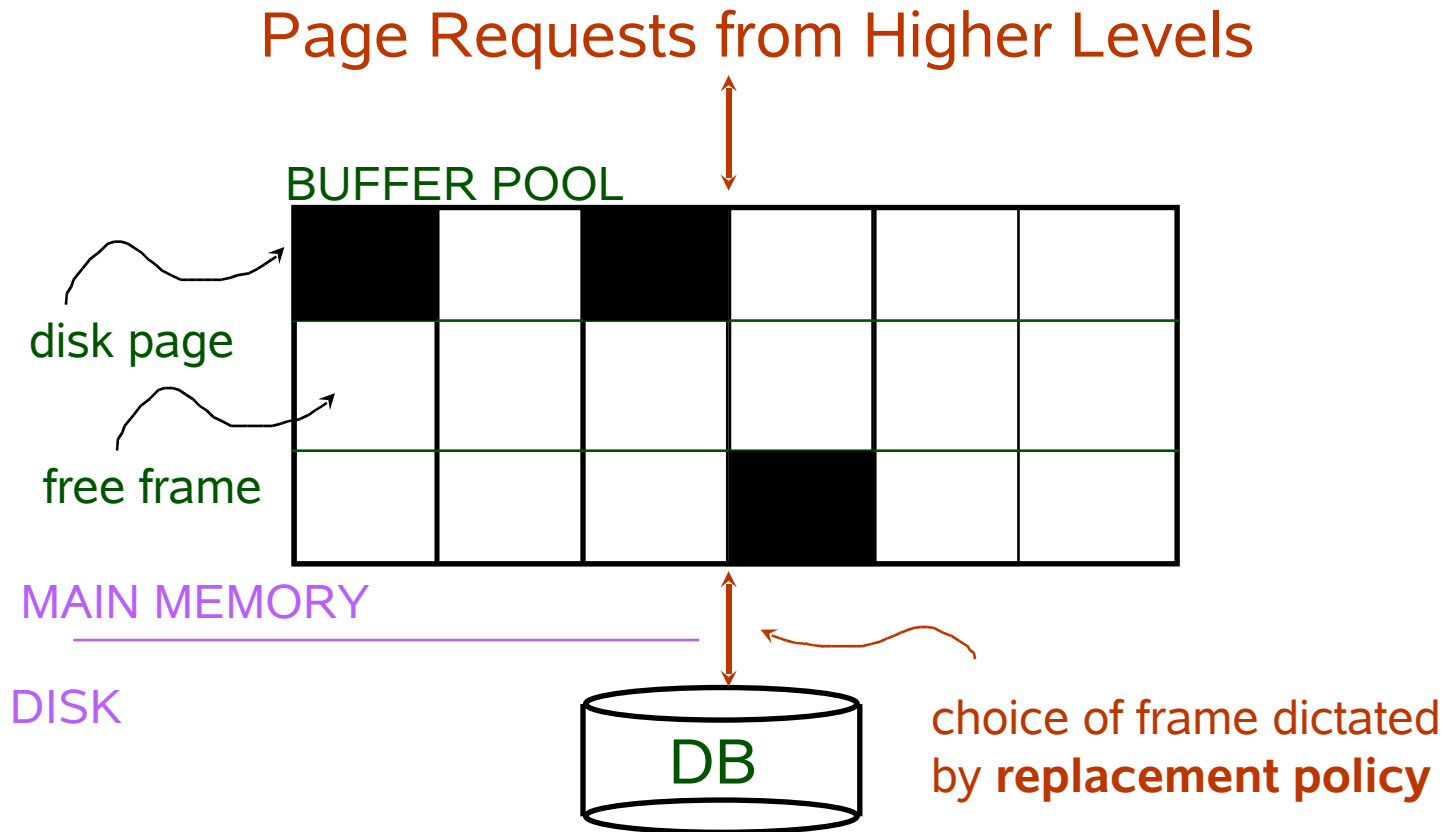
Outline

- ❖ The memory hierarchy
- ❖ RAID
- ❖ Disk space management
- ❖ Files of records
- ❖ Page formats
- ❖ Record formats

Buffer Manager

- ❖ Example:
 - DB contains 1 million pages
 - 100 pages can be held in main memory
 - Query: scan 1 million pages.
- ❖ Since all the data can not be brought to main memory, DBMS
 - brings pages to main memory when they are needed.
 - Decide existing page to replace when there is no space
- ❖ Buffer manager is a software layer responsible to bring the pages to main memory.
- ❖ Higher-level code can be written whether data pages are in memory or not.

Buffer Management in a DBMS



- ❖ *Data must be in RAM for DBMS to operate on it!*
- ❖ *Table of $\langle \text{frame\#}, \text{pageid} \rangle$ pairs is maintained.*

When a Page is Requested ...

- ❖ If requested page is not in pool:
 - Choose a frame for *replacement*
 - If frame is dirty, write it to disk
 - Read requested page into chosen frame
 - ❖ *Pin* the page and return its address.
- ➡ *If requests can be predicted (e.g., sequential scans) pages can be pre-fetched several pages at a time!*

More on Buffer Management

- ❖ Requestor of page must unpin it, and indicate whether page has been modified:
 - *dirty* bit is used for this.
- ❖ Page in pool may be requested many times,
 - a *pin count* is used. A page is a candidate for replacement iff $pin\ count = 0$.
- ❖ CC & recovery may entail additional I/O when a frame is chosen for replacement. (*Write-Ahead Log* protocol; more later.)

Buffer Replacement Policy

- ❖ Frame is chosen for replacement by a *replacement policy*:
 - Least-recently-used (LRU), Clock, MRU etc.
- ❖ Policy can have big impact on # of I/O's; depends on the *access pattern*.
- ❖ *Sequential flooding*: Nasty situation caused by LRU + repeated sequential scans.
 - # buffer frames < # pages in file means each page request causes an I/O. MRU much better in this situation (but not in all situations, of course).

DBMS vs. OS File System

OS does disk space & buffer mgmt: why not let OS manage these tasks?

- ❖ Differences in OS support: portability issues
- ❖ Some limitations, e.g., files can't span disks.
- ❖ Buffer management in DBMS requires ability to:
 - **pin a page** in buffer pool, **force a page** to disk (important for implementing CC & recovery),
 - adjust *replacement policy*, and **pre-fetch pages** based on access patterns in typical DB operations.

Outline

- ❖ The memory hierarchy
- ❖ RAID
- ❖ Disk space management
- ❖ Buffer manager
- ❖ Files of records
- ❖ Page formats
- ❖ Record formats

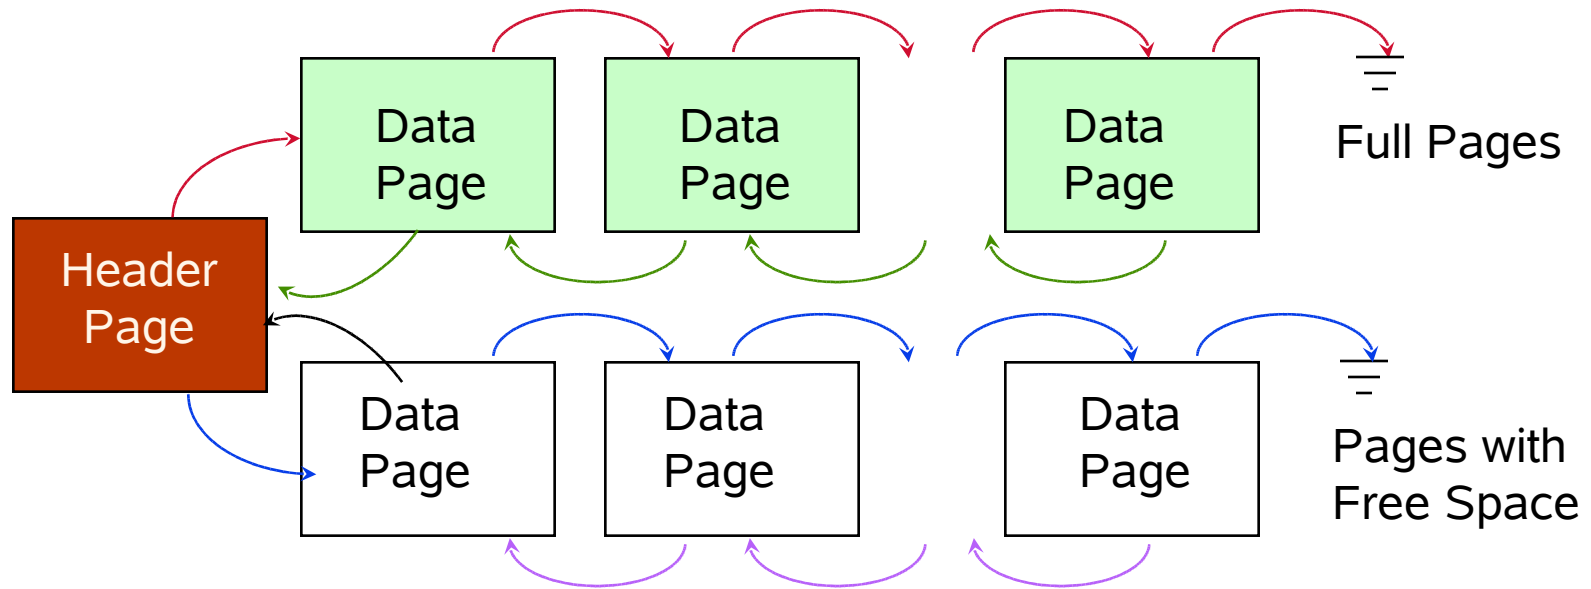
Files of Records

- ❖ Page or block is OK when doing I/O, but higher levels of DBMS operate on *records*, and *files of records*.
- ❖ FILE: A collection of pages, each containing a collection of records. Must support:
 - insert/delete/modify record
 - read a particular record (specified using *record id*)
 - scan all records (possibly with some conditions on the records to be retrieved)

Unordered (Heap) Files


- ❖ Simplest file structure contains records in no particular order.
- ❖ As file grows and shrinks, disk pages are allocated and de-allocated.
- ❖ To support record level operations, we must:
 - keep track of the *pages* in a file
 - keep track of *free space* on pages
 - keep track of the *records* on a page
- ❖ There are many alternatives for keeping track of this.

Heap File Implemented as a List



- ❖ The header page id and Heap file name must be stored someplace.
- ❖ Each page contains 2 'pointers' plus data.
- ❖ Disadvantage:
 - All pages in a file is in free list if records are of variable length.

Heap File Using a Page Directory



- ❖ The entry for a page can include the number of free bytes on the page.
- ❖ The directory is a collection of pages; linked list implementation is just one alternative.
 - *Much smaller than linked list of all HF pages!*


Page Formats

- ❖ Page abstraction is appropriate when dealing with I/O issues.
- ❖ Higher levels of DBMS see data as a collection of records.
- ❖ A page is a collection of slots, where each slot contains a record.
- ❖ It can be identified with $\langle \text{page-id}, \text{slot-id} \rangle$, which is a record id.

Fixed length records

- ❖ If the records of a page are guaranteed to be of the same length, records can be arranged consecutively within the page.
- ❖ At any time, some slots are occupied and others are unoccupied.
- ❖ When a record is inserted, we must locate empty slot and place the record.
- ❖ Issue: How to keep track of empty slots and how to locate all records on the page ?
- ❖ Alternative 1: Store records in the first N slots, when the record is deleted, move the last record in the page to the vacated slot.
 - It will not work, if there is external reference.
- ❖ Alternative 2: Handle deletions using array of bits, one per slot, to keep track of free slot information.

Page Formats: Fixed Length Records




➡ Record id = $\langle \text{page id}, \text{slot \#} \rangle$. In first alternative, moving records for free space management changes rid; may not be acceptable.

Variable-length records


- ❖ If records are of variable length, we can not divide the page into a fixed collection of records.
 - Problem: finding the empty slot of right length
 - If the slot is too big, it wastes space!
- ❖ Solution: Maintain a directory of slots
 - Record offset and record length for each record.
 - Record offset is the pointer to the record
- ❖ The free space available should be managed carefully, as the page is not preformatted into slots.

Page Formats: Variable Length Records



- ➡ *Can move records on page without changing rid; so, attractive for fixed-length records too.*


Record Formats: Fixed Length



- ❖ Information about field types same for all records in a file; stored in *system catalogs*.
- ❖ Finding *i*'th field requires scan of record.

Record Formats: Variable Length

- ❖ Two alternative formats (# fields is fixed):



- ➡ Second offers direct access to i'th field, efficient storage of nulls (special *don't know* value); small directory overhead.

Outline

- ❖ The memory hierarchy
- ❖ RAID
- ❖ Disk space management
- ❖ Files of records
- ❖ Page formats
- ❖ Record formats

System Catalogs

- ❖ For each index:
 - structure (e.g., B+ tree) and search key fields
 - ❖ For each relation:
 - name, file name, file structure (e.g., Heap file)
 - attribute name and type, for each attribute
 - index name, for each index
 - integrity constraints
 - ❖ For each view:
 - view name and definition
 - ❖ Plus statistics, authorization, buffer pool size, etc.
- ➡ *Catalogs are themselves stored as relations!*

Attr_Cat(attr_name, rel_name, type, position)

attr_name	rel_name	type	position
attr_name	Attribute_Cat	string	1
rel_name	Attribute_Cat	string	2
type	Attribute_Cat	string	3
position	Attribute_Cat	integer	4
sid	Students	string	1
name	Students	string	2
login	Students	string	3
age	Students	integer	4
gpa	Students	real	5
fid	Faculty	string	1
fname	Faculty	string	2
sal	Faculty	real	3

Summary

- ❖ Disks provide cheap, non-volatile storage.
 - Random access, but cost depends on location of page on disk; important to arrange data sequentially to minimize *seek* and *rotation* delays.
- ❖ Buffer manager brings pages into RAM.
 - Page stays in RAM until released by requestor.
 - Written to disk when frame chosen for replacement (which is sometime after requestor releases the page).
 - Choice of frame to replace based on *replacement policy*.
 - Tries to *pre-fetch* several pages at a time.

Summary (Contd.)

❖ DBMS vs. OS File Support

- DBMS needs features not found in many OS's, e.g., forcing a page to disk, controlling the order of page writes to disk, files spanning disks, ability to control pre-fetching and page replacement policy based on predictable access patterns, etc.
- ❖ Variable length record format with field offset directory offers support for direct access to i'th field and null values.
- ❖ Slotted page format supports variable length records and allows records to move on page.

Summary (Contd.)

- ❖ File layer keeps track of pages in a file, and supports abstraction of a collection of records.
 - Pages with free space identified using linked list or directory structure (similar to how pages in file are kept track of).
- ❖ Indexes support efficient retrieval of records based on the values in some fields.
- ❖ Catalog relations store information about relations, indexes and views. (*Information that is common to all records in a given collection.*)