

# Hash-Based Indexes


## Chapter 10

# Introduction

- ❖ *As for any index, 3 alternatives for data entries  $k^*$ :*
  - Data record with key value  $k$
  - $\langle k, \text{rid of data record with search key value } k \rangle$
  - $\langle k, \text{list of rids of data records with search key } k \rangle$
  - Choice orthogonal to the *indexing technique*
- ❖ Hash-based indexes are best for *equality selections*. **Cannot** support range searches.
- ❖ Static and dynamic hashing techniques exist; trade-offs similar to ISAM vs. B+ trees.

# Static Hashing

- ❖ # primary pages fixed, allocated sequentially, never de-allocated; overflow pages if needed.
- ❖  $h(k) \bmod M = \text{bucket to which data entry with key } k \text{ belongs. (} M = \# \text{ of buckets)}$



# Static Hashing (Contd.)

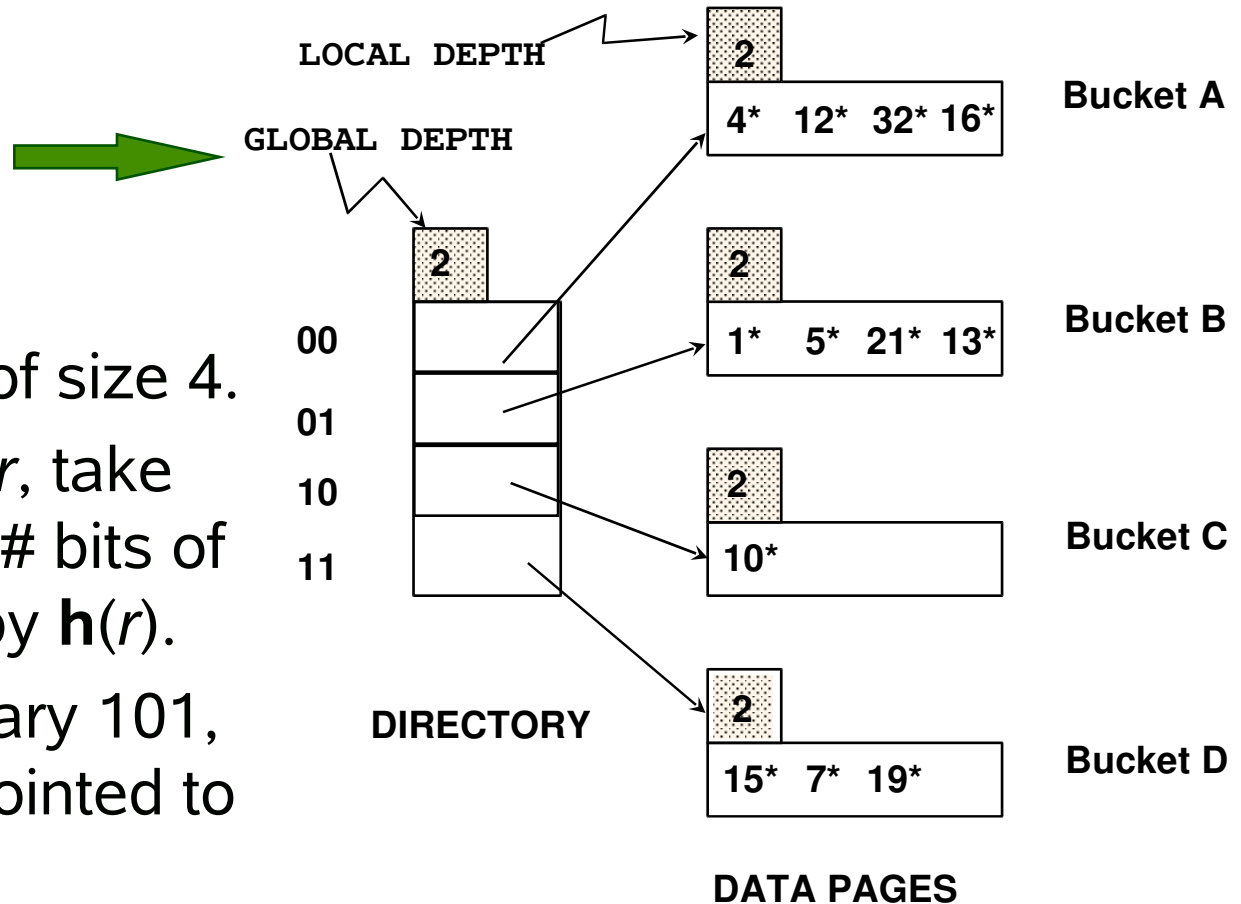
- ❖ Buckets contain *data entries*.
- ❖ Hash fn works on *search key* field of record *r*.  
Must distribute values over range 0 ... M-1.
  - $h(key) = (a * key + b)$  usually works well.
  - a and b are constants; lots known about how to tune **h**.
- ❖ **Long overflow chains** can develop and degrade performance.
  - **Extendible** and **Linear Hashing**: Dynamic techniques to fix this problem.

# Extendible Hashing

- ❖ Objective is to avoid overflow buckets
- ❖ Situation: Bucket (primary page) becomes full. Instead of overflow bucket, why not re-organize file by *doubling* # of buckets?
  - Reading and writing all pages is expensive!
  - Idea: Use directory of pointers to buckets, double # of buckets by *doubling the directory*, splitting just the bucket that overflowed!
  - Directory much smaller than file, so doubling it is much cheaper. Only one page of data entries is split. *No overflow page!*
  - Trick lies in how hash function is adjusted!

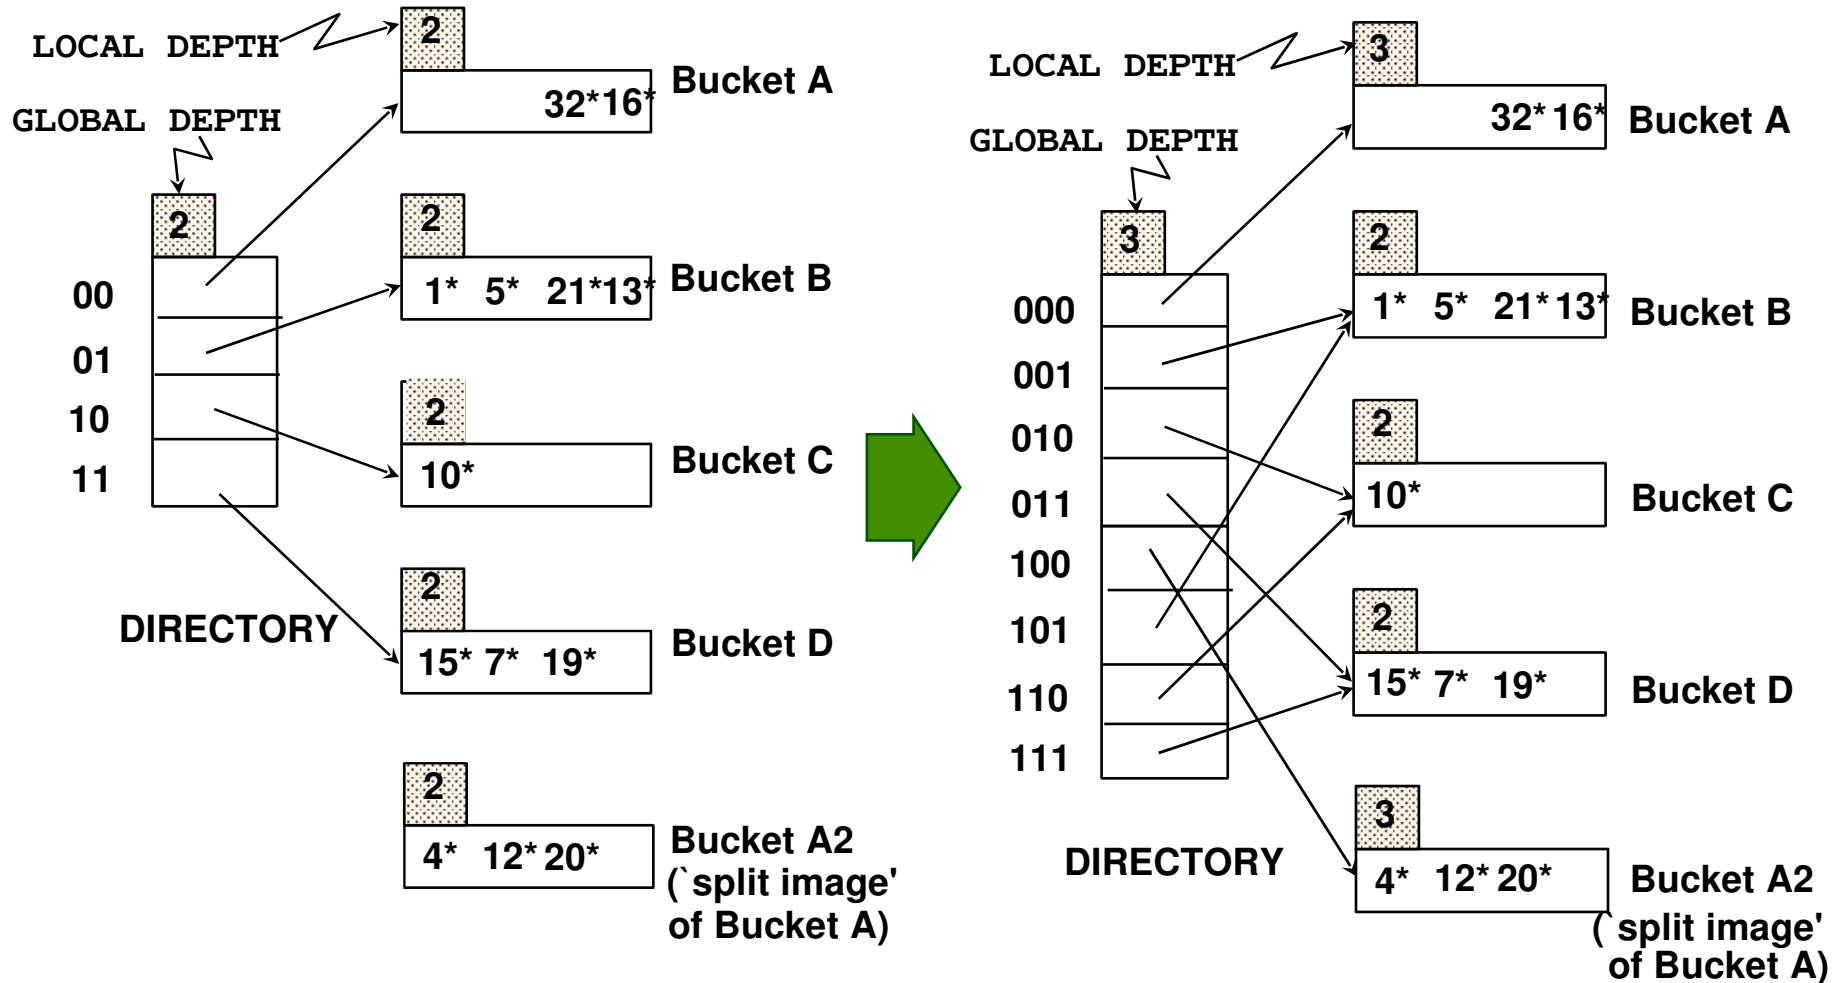
# Example

- ❖ Directory is array of size 4.
- ❖ To find bucket for  $r$ , take last '*global depth*' # bits of  $\mathbf{h}(r)$ ; we denote  $r$  by  $\mathbf{h}(r)$ .
  - If  $\mathbf{h}(r) = 5 = \text{binary } 101$ , it is in bucket pointed to by 01.



- ❖ **Insert:** If bucket is full, *split* it (*allocate new page, re-distribute*).
- ❖ *If necessary*, double the directory. (As we will see, splitting a bucket does not always require doubling; we can tell by comparing *global depth* with *local depth* for the split bucket.)

# Insert $h(r)=20$ (Causes Doubling)



# Points to Note


- ❖ 20 = binary 10100. Last **2** bits (00) tell us  $r$  belongs in A or A2. Last **3** bits needed to tell which.
  - *Global depth of directory*: Max # of bits needed to tell which bucket an entry belongs to.
  - *Local depth of a bucket*: # of bits used to determine if an entry belongs to this bucket.
- ❖ When does bucket split cause directory doubling?
  - Before insert, *local depth* of bucket = *global depth*. Insert causes *local depth* to become  $>$  *global depth*; directory is doubled by *copying it over* and 'fixing' pointer to split image page. (Use of least significant bits enables efficient doubling via copying of directory!)



# Directory Doubling


Why use least significant bits in directory?

⇒ Allows for doubling via copying!



Least Significant

vs.



Most Significant

# Comments on Extendible Hashing

- ❖ If directory fits in memory, equality search answered with one disk access; else two.
  - 100MB file, 100 bytes/rec, 4K pages contains 1,000,000 records (as data entries) and 25,000 directory elements; chances are high that directory will fit in memory.
  - Directory grows in spurts, and, if the distribution of *hash values* is skewed, directory can grow large.
  - Multiple entries with same hash value cause problems!
- ❖ **Delete**: If removal of data entry makes bucket empty, can be merged with 'split image'. If each directory element points to same bucket as its split image, can halve directory.

# Linear Hashing


- ❖ This is another dynamic hashing scheme, an alternative to Extendible Hashing.
- ❖ LH handles the problem of long overflow chains without using a directory, and handles duplicates.
- ❖ Idea: Use a family of hash functions  $h_0, h_1, h_2, \dots$ 
  - $h_i(\text{key}) = h(\text{key}) \bmod(2^i N)$ ;  $N$  = initial # buckets
  - $h$  is some hash function (range is *not* 0 to  $N-1$ )
  - If  $N = 2^{d_0}$ , for some  $d_0$ ,  $h_i$  consists of applying  $h$  and looking at the last  $d_i$  bits, where  $d_i = d_0 + i$ .
  - $h_{i+1}$  doubles the range of  $h_i$  (similar to directory doubling)
  - If  $h_i$  maps data entry to one of  $M$  buckets and  $h_{i+1}$  maps to one of  $2M$  buckets.

# Linear Hashing (Contd.)

- ❖ Directory avoided in LH by using overflow pages, and choosing bucket to split round-robin.
  - **Splitting proceeds in `rounds`.** Round ends when all  $N_R$  initial (for round  $R$ ) buckets are split. Buckets 0 to ***Next-1*** have been split; *Next* to  $N_R$  yet to be split.
  - **Current round number is *Level*.**
  - **Search:** To find bucket for data entry  $r$ , find  **$h_{Level}(r)$** :
    - If  $h_{Level}(r)$  in range `*Next* to  $N_R$ ',  $r$  belongs here.
    - Else,  $r$  could belong to bucket  $h_{Level}(r)$  or bucket  $h_{Level}(r) + N_R$ ; must apply  $h_{Level+1}(r)$  to find out.

# Overview of LH File

❖ In the middle of a round.



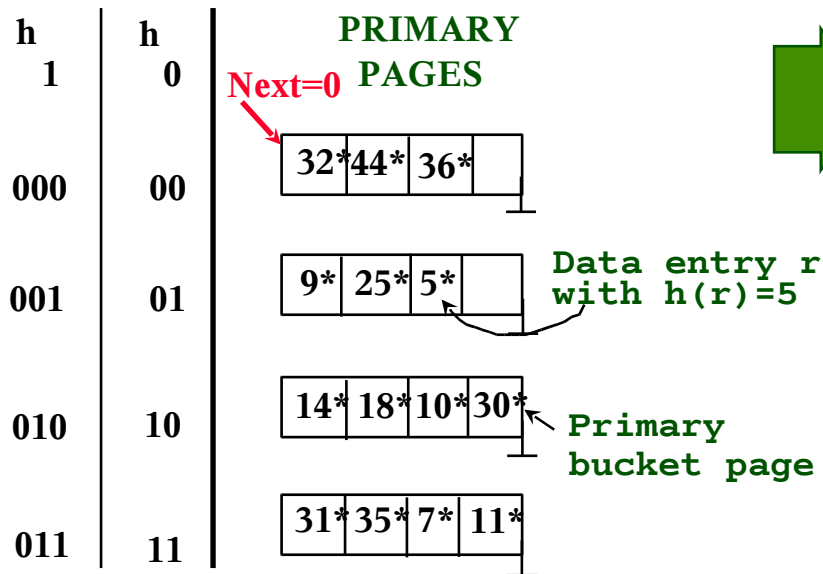
# Linear Hashing (Contd.)

- ❖ **Insert:** Find bucket by applying  $h_{Level} / h_{Level+1}$ :
  - If bucket to insert into is full:
    - Add overflow page and insert data entry.
    - (*Maybe*) Split *Next* bucket and increment *Next*.
- ❖ Can choose any criterion to `trigger' split.
- ❖ Since buckets are split round-robin, long overflow chains don't develop!
- ❖ Doubling of directory in Extendible Hashing is similar; switching of hash functions is *implicit* in how the # of bits examined is increased.

# Example of Linear Hashing

- ❖ On **split**,  $h_{\text{Level}+1}$  is used to **re-distribute** entries.


Level=0, N=4



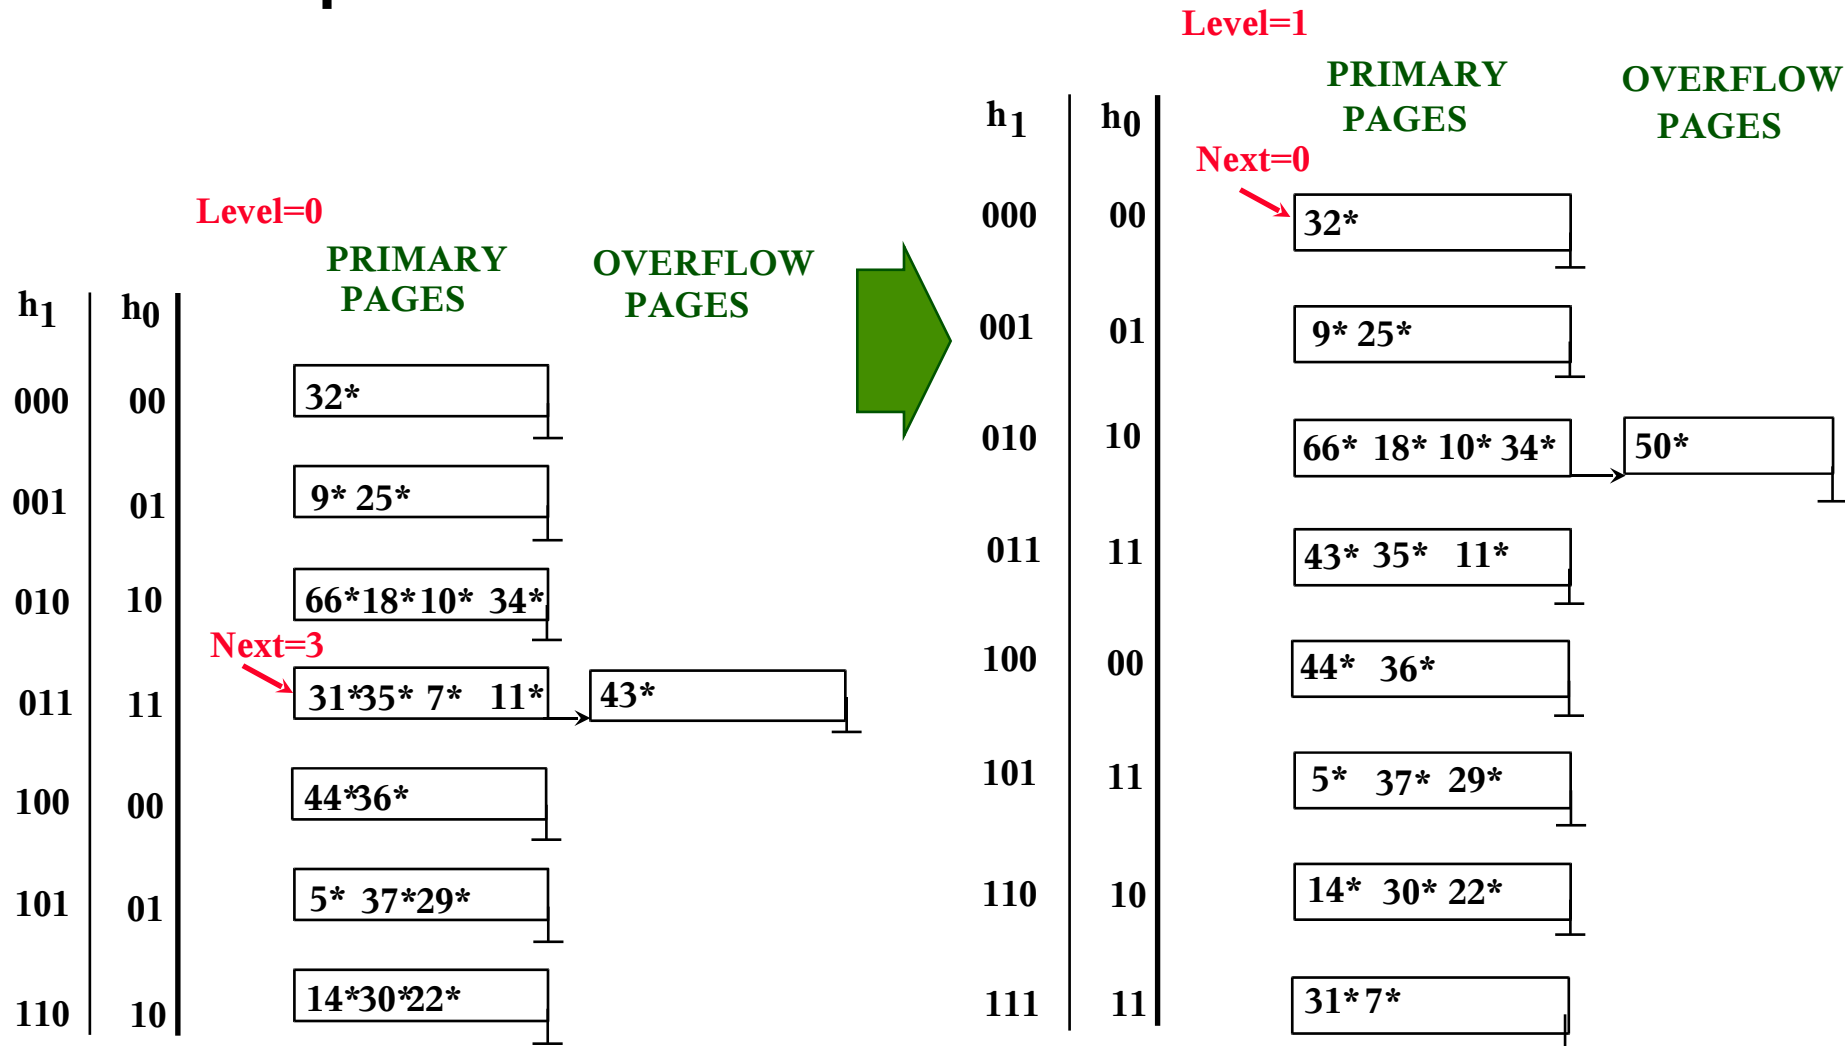
*(This info is for illustration only!)*

*(The actual contents of the linear hashed file)*

Level=0



# Example: End of a Round





# LH Described as a Variant of EH

- ❖ The two schemes are actually quite similar:
  - Begin with an EH index where directory has  $N$  elements.
  - Use overflow pages, split buckets round-robin.
  - First split is at bucket 0. (Imagine directory being doubled at this point.) But elements  $\langle 1, N+1 \rangle$ ,  $\langle 2, N+2 \rangle$ , ... are the same. So, need only create directory element  $N$ , which differs from 0, now.
    - When bucket 1 splits, create directory element  $N+1$ , etc.
- ❖ So, directory can double gradually. Also, primary bucket pages are created in order. If they are *allocated* in sequence too (so that finding  $i$ 'th is easy), we actually don't need a directory!

# Summary

- ❖ Hash-based indexes: best for equality searches, cannot support range searches.
- ❖ Static Hashing can lead to long overflow chains.
- ❖ Extendible Hashing avoids overflow pages by splitting a full bucket when a new data entry is to be added to it.
  - Directory to keep track of buckets, doubles periodically.
  - Can get large with skewed data; additional I/O if this does not fit in main memory.

# Summary (Contd.)

- ❖ Linear Hashing avoids directory by splitting buckets round-robin, and using overflow pages.
  - Overflow pages not likely to be long.
  - Duplicates handled easily.
  - Space utilization could be lower than Extendible Hashing, since splits not concentrated on `dense' data areas.
    - Can tune criterion for triggering splits to trade-off slightly longer chains for better space utilization.
- ❖ For hash-based indexes, a *skewed* data distribution is one in which the *hash values* of data entries are not uniformly distributed!