

Transaction Management Overview

Chapter 16

Unsophisticated users (Customers,
Travel agents)

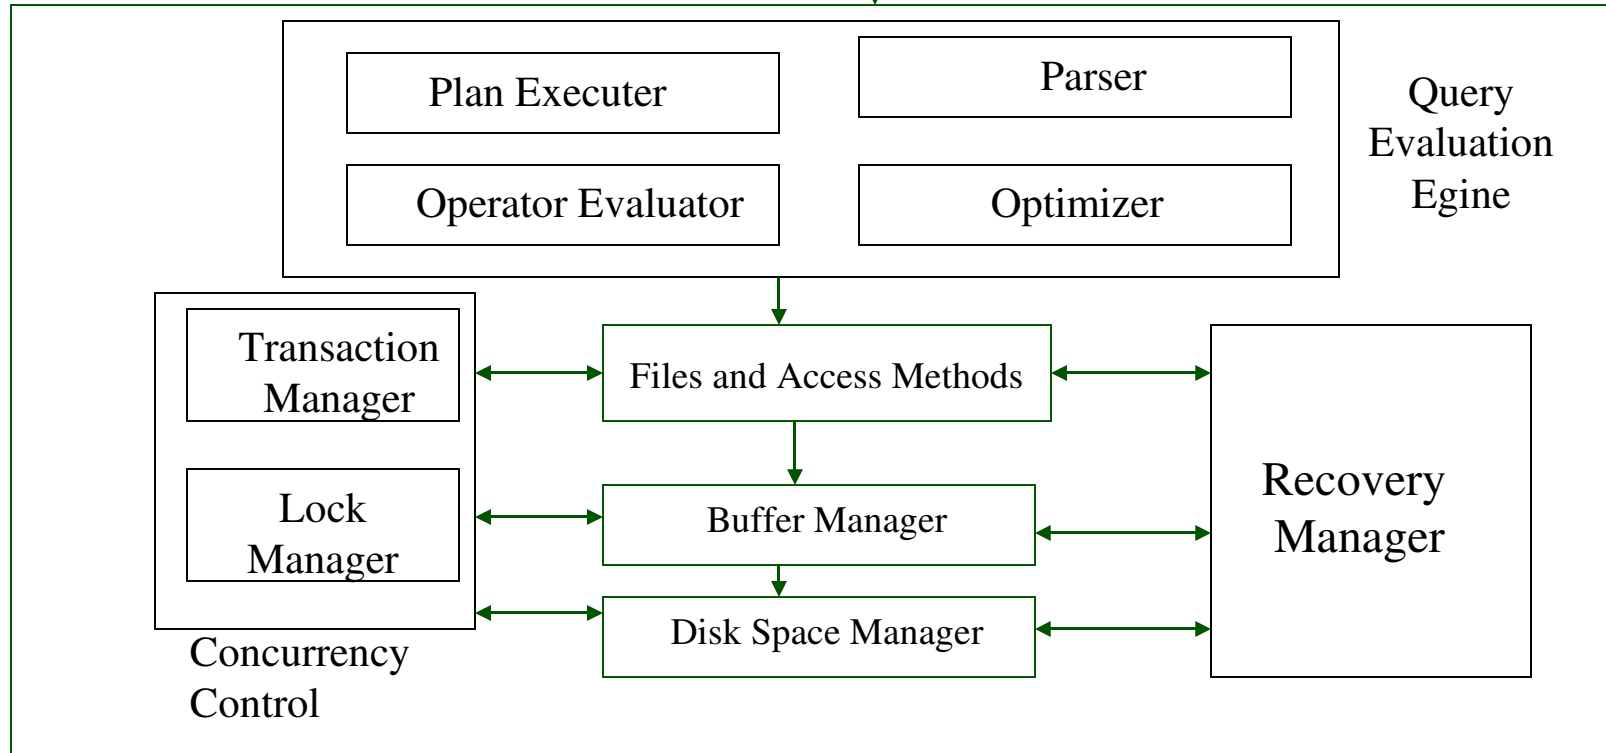
Sophisticated users, application
Programmers, DB Administrators

Web Forms


Application Front Ends

SQL Interface

SQL COMMANDS



Architecture
Of
DBMS



Outline

- ❖ The ACID Properties
- ❖ Transactions and Schedules
- ❖ Concurrent Execution of Transactions
- ❖ Lock-based Concurrency control
- ❖ Performance of Locking
- ❖ Transaction Support in SQL
- ❖ Introduction to Crash Recovery

Concurrency

- ❖ **Concurrent** execution of user programs is essential for good DBMS performance.
- ❖ Because disk accesses are frequent, and relatively slow, it is important to keep the CPU humming by working on multiple user programs concurrently.

Transactions

- ❖ A user's program may carry out many operations on the data from the database,
- ❖ but the DBMS is only concerned about what data is **read/written from/to** the database.
- ❖ A **transaction** is the DBMS's abstract view of a user program: a sequence of **reads and writes**

Example of User Program and Transaction

- ❖ Assume **A** is a tuple in relation **Account**
- ❖ User program

```
Calculate_Interests (Account A, Rate r) {  
    float x,y;  
    x := A; // read from database  
    if (x <= 1000) y := 1.05*x  
    else if (x <= 10000) y := 1.06*x  
    else y := 1.065*x;  
    A := y; // write to database  
}
```

- ## ❖ Corresponding Transaction T

$T: x := A, A := y$ or $T: R(A), W(A)$

Transactions

- ❖ A transaction is a sequence of operations.
- ❖ The operations can be reading $R(A)$, writing $W(A)$, commit or abort
- ❖ Required Properties of transactions:
 - Atomicity
 - Consistency
 - Isolation
 - Durability
- ❖ Called as “ACID properties”

Atomicity of Transactions

- ❖ A transaction might **commit** after all its actions, or it could **abort** (or be aborted) after executing some actions.
- ❖ when it commits, all its writes are written
when it aborts, none of its writes is written

This is called the **atomicity** property

Atomicity

- ❖ Why atomicity?
- ❖ Partial effects leave wrong data in DB

Transfer (Account A, Account B, Money r)

{ $A := A - r$; $B := B + r$ } //Transfer \$r from A
to B

- ❖ What if **Transfer** is aborted after $A := A - r$
and before $B := B + r$?
- DBMS **logs** all actions so that it can **undo** the
actions of aborted transactions.
- Transaction Manager and RECOVERY

Consistency

- ❖ Each transaction must leave the database in a consistent state if the DB is consistent when the transaction begins (called **consistency** property).
- ❖ DBMS assumes that the consistency holds for each transaction
- ❖ Ensuring consistency property is the responsibility of the user.

Isolation

- ❖ Users submit transactions, and can think of each transaction as executing by itself (called **isolation** property).
 - concurrently executing transactions will have the same effect as “one at a time”
 - Transactions are isolated or protected.
 - Isolation is achieved by the DBMS

Isolation

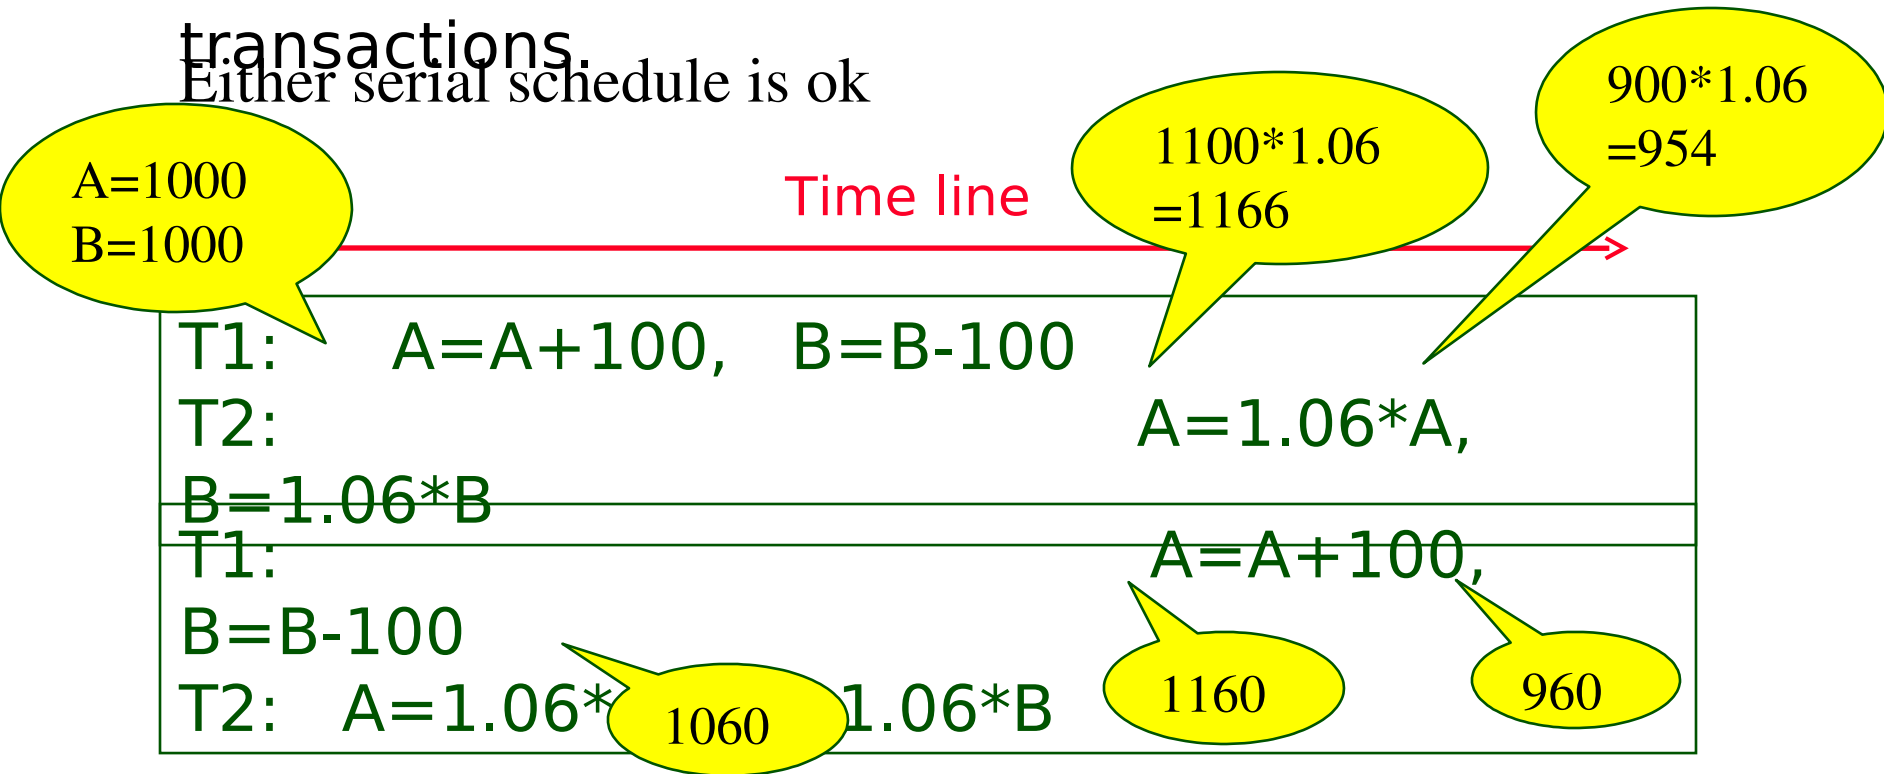
- ❖ Consider two transactions :

T1:	BEGIN	A=A+100,	B=B-100	END
T2:	BEGIN	A=1.06*A,	B=1.06*B	END

- ❖ T1 transfers \$100 from B's account to A's account.
- ❖ T2 credits both accounts with a 6% interest payment.
- ❖ There is no guarantee that T1 will execute before T2 or vice-versa, if both are submitted together.
- ❖ However, the net effect *must* be equivalent to these two transactions running **serially** in some order.

Serial schedule: Schedule that does not interleave the actions of different transactions.

Either serial schedule is ok



A non-serial execution (interleaving schedule):

T1:	$A=A+100$,	$B=B-100$
T2:	$A=1.06*A$,	
	$B=1.06*B$	

Durability

- ❖ Once a transaction commits, its effect persists (**durability** property.)
- ❖ ACID: Atomicity, Consistency, Isolation, Durability

Outline

- ❖ The ACID Properties
- ❖ **Transactions and Schedules**
- ❖ Concurrent Execution of Transactions
- ❖ Lock-based Concurrency control
- ❖ Performance of Locking
- ❖ Transaction Support in SQL
- ❖ Introduction to Crash Recovery

A Schedule

- ❖ DBMS sees transaction as a list of actions: reads, writes
- ❖ $R_T(O)$: Reading object O .
- ❖ $W_T(O)$: writing object O .
- ❖ In addition, commit and abort actions are included.
 - $Commit_T$ denotes committing action
 - $Abort_T$ denotes aborting action.
- ❖ $T = \{T_1, \dots, T_n\}$ is a set of transactions.
- ❖ A **schedule (history)** is a sequence of operations in T_1, \dots, T_n such that for each
 1. Each operation in T_i appears exactly once, and

Transaction Interleaving

Consider a possible interleaving **schedule**:

A=1000
B=1000

Time line

$900 * 1.06$
 $= 954$

T1:	A=A+100,	B=B-100
T2:	A=1.06*A,	B=1.06*B

❖ This is OK. It is equivalent to the following serial execution

T1:	A=A+100, B=B-100
T2:	

A=1.06*A, B=1.06*B


$1100 * 1.06$
 $= 1166$

$900 * 1.06$
 $= 954$

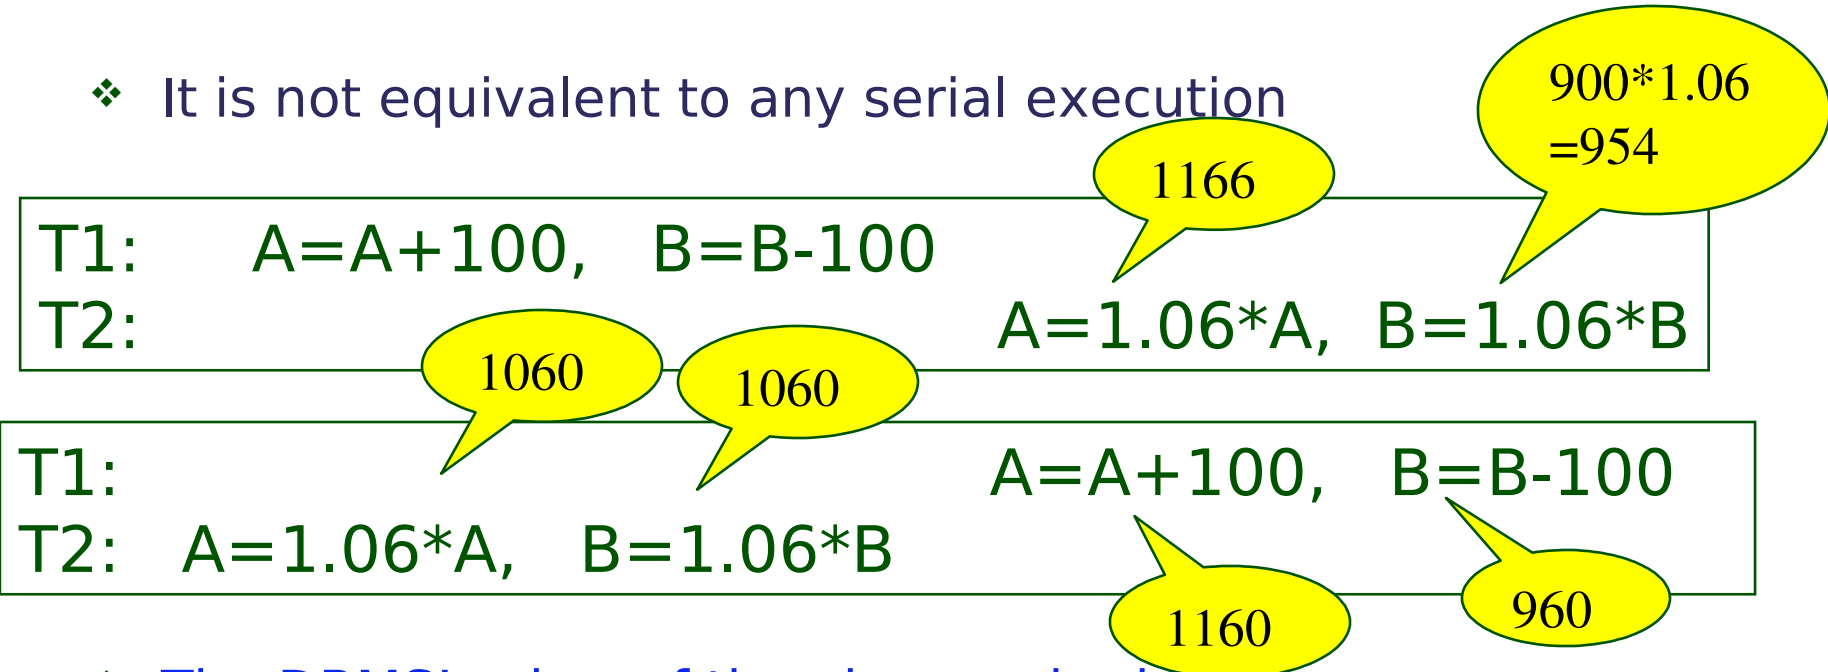
❖ But what about:

T1:	A=A+100,	B=B-100
T2:	A=1.06*A, B=1.06*B	


❖ But what about:



❖ It is not equivalent to any serial execution



❖ The DBMS's view of the above schedule:



Scheduling Transactions

T1: $A=A+100$, $B=B-100$

~~T2:~~ $A=1.06*A$, $B=1.06*B$

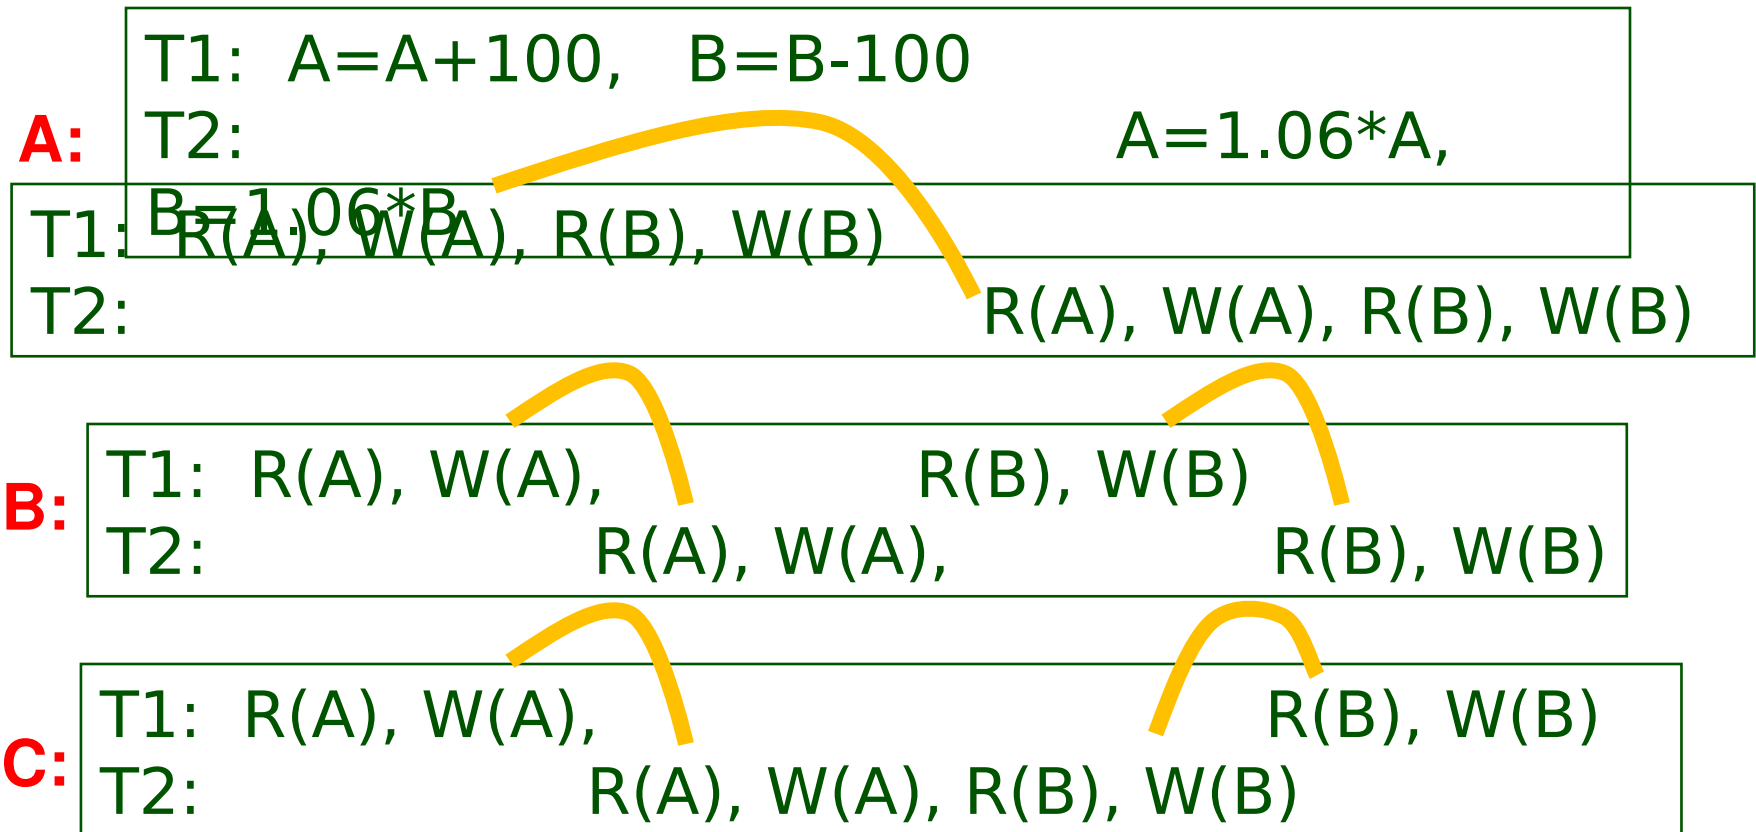
- ❖ **Equivalent schedules:** the effects (on the database) of two schedules are identical

B: T1: $A=A+100$, $B=B-100$
T2: $A=1.06*A$, $B=1.06*B$

C: T1: $A=A+100$, $B=B-100$
T2: $A=1.06*A$, $B=1.06*B$

- ❖ A and B are equivalent, but A and C are not.

❖ Translate this example to DBMS view, we have:



❖ A and B are equivalent, but A and C are not.

❖ Problem: **WR conflict**

Serializable Schedule

- ❖ A schedule that is equivalent to some serial execution of the transactions is called a **serializable schedule**.
- ❖ every serializable schedule preserves consistency
 - In the serial schedule, transactions are executed **one after another**
 - Every transaction itself preserves consistency (by consistency property)
 - So the serial schedule preserves consistency
 - So the serializable schedule preserves consistency

Anomalies with Interleaved Execution

- ❖ Reading Uncommitted Data
- ❖ (WR Conflicts, “dirty reads”)

Should not have
allowed T2
to withdraw
\$500 !!!

Comm
it T2 !

T1:	R(A), W(A),	Abort
T2:	R(A), W(A), C	

Initially
A=0

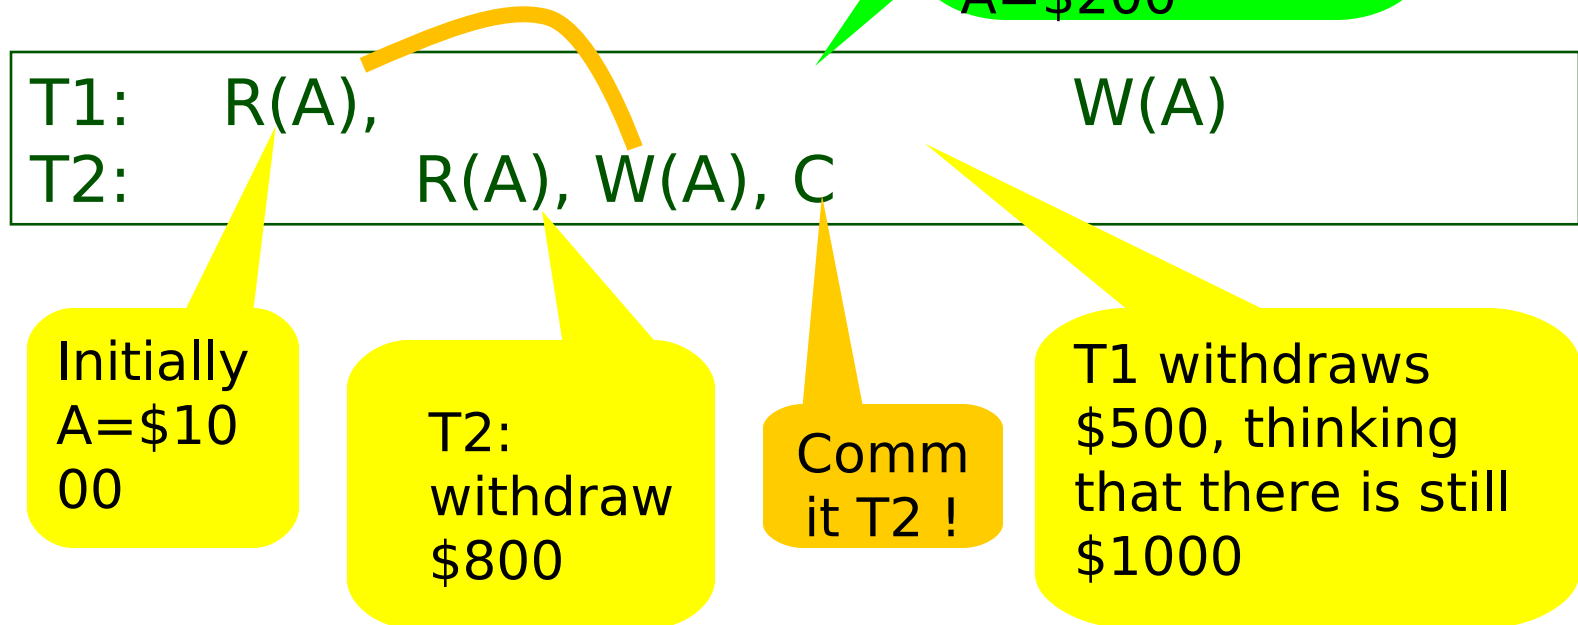
T1 :
Deposit
\$1000

T2 :
Withdra
w
\$500

T1 decides not to
deposit
the \$1000 after
all

Anomalies with Interleaved Execution


- ❖ Unrepeatable Reads
- ❖ (RW Conflicts):



- ❖ Imagine T1 verifies first that fund is sufficient, then withdraw.

Anomalies

❖ Overwriting Uncommitted Data (WW Conflicts):



- Both A and B must be initialized to 1 or 0.
- T1 writes 1 to both, and T2 writes 0 to both.
- In the end, A = 0, B = 1
(T1 writes 1 to A, but overwritten by T2)
Now A and B have got different values

None of these anomalies is a serializable schedule!!

Recoverable Schedules

- ❖ In a recoverable schedule, transactions commit only after all transactions whose changes they read commit.
 - Cascading abort problem
- ❖ The schedule becomes recoverable, if a transactions read the effect of committed transactions.
- ❖ The following schedule is not recoverable

T1:	R(A) W(A),	A
T2:	R(A) W(A) R(B) W(B)	C

Concurrency Control

- ❖ DBMS must
 - Ensure serializable schedules
 - Recoverable schedules
 - Improve parallelism for high throughput and low response time.
- ❖ A DBMS uses a locking-based protocol.
- ❖ Strict Two-phase Locking (Strict 2PL) Protocol:
 - transaction must obtain a **S (shared) lock** on object before reading, and an **X (exclusive) lock** on object before writing.
 - If a transaction holds an X lock on an object, no other Xact can get a lock (S or X) on that object.
 - All locks held by a transaction are released when the transaction completes

Strict 2PL allows only serializable schedules.

S-X lock conflict and X-X lock conflict

❖ T1: R(A), W(A),	Abort
❖ T2:	R(A), W(A), C

A is X-locked by T1 !

T1: R(A),	R(A), W(A),
Abort	
T2:	R(A), W(A), C

A is S-locked by T1 !

T1: W(A),	W(B), C
T2:	W(A), W(B), C

A is X-locked by T1 !

Aborting a Transaction

- ❖ If a transaction T_i is aborted, all its actions have to be undone.
if T_j reads an object last written by T_i , T_j must be aborted as well!
- ❖ Most systems avoid such **cascading aborts** by releasing a transaction's locks only at commit time.
 - If T_i writes an object, T_j can read this only after T_i commits.

Aborting a Transaction

- ❖ In order to undo the actions of an aborted transaction, the DBMS maintains a **log** in which every write is recorded.
- ❖ This mechanism is also used to recover from system crashes:
all active transactions at the time of the crash are aborted when the system comes back up.

Deadlocks

- ❖ A cycle of transactions waiting for a lock
- ❖ Timeout is a simple mechanism
- ❖ Details will be discussed later.

Performance of Locking

- ❖ Lock-based schemes resolve conflicts using two basic mechanisms: blocking and aborting.
 - Both involve performance penalty
- ❖ Blocking leads to cascading waiting
 - CPU is free but can not be used!
 - Deadlocks- extreme blocking!
 - About 1 % of transactions are involved in deadlock.
 - Transaction blocking is significant overhead
- ❖ Thrashing occurs as MPL increases
 - Occurs if about 30% of transactions are blocked.
- ❖ Throughput can be increased
 - By locking the small sized objects possible.
 - By reducing the time that transaction holds
 - By reducing the **hotspots**
 - A hotspot is a data object that is frequently accessed and modified.

Transaction Support in SQL

- ❖ So far we have studied abstract model
- ❖ A transaction automatically starts when user executes a statement that accesses the database or catalogs: SELECT, UPDATE, CREATE TABLE
- ❖ The statements are executed until COMMIT and ROLLBACK (for abort) commands is encountered.
- ❖ In SQL:1999 the following features are provided
 - Savepoint: allows us to identify a point and selectively rollback operations carried out after this point.
 - SAVEPOINT (Name)
 - ROLLBACK TO SAVEPOINT (Savepoint name)

What should we lock ?

- ❖ Transaction locks objects
- ❖ Consider the following query
 - `SELECT S.rating, MIN (S.age)`
 - `FROM Sailors S`
 - `WHERE S.rating=8`
- ❖ Option 1:
 - DBMS should put a shared lock on entire Table. Another Transaction which modifies Sailor table should have exclusive lock on the table. LEADS to low concurrency.
 - We can lock the smaller objects
 - Set an exclusive lock on single row.
 - Reading can be carried out in parallel
- ❖ So, DBMS can lock the objects at different granularities!
 - Tables-level and row-level

Phantom Problem

- ❖ Suppose DBMS sets row-level locks with rating =8 for T1.
- ❖ T3 can add another row and lock it.
- ❖ Then T1 gets two results based on execution of T1 relative to T2.
- ❖ Phantom problem: A transaction retrieves collection of objects twice and sees different results, even though it does not modify any results.
- ❖ Solution: lock entire table.

Transaction characteristics in SQL

- ❖ SQL allows the programmers to specify three characteristics of transaction.
 - Access mode
 - Diagnostics size
 - Isolation level
- ❖ Diagnostics size:
 - determines the number of error conditions that can be recorded.
- ❖ Access mode:
 - if access mode is READ ONLY, the transaction is not allowed to modify the database.
 - If the mode is READ-WRITE
 - INSERT, DELETE, UPDATE and CREATE commands can be executed.
- ❖ Isolation Level:
 - controls the extent to which a given transaction is exposed to the actions of other transactions executing concurrently.
 - By choosing one of four possible isolation level settings, the user can obtain greater concurrency.
- ❖ Isolation levels: READ UNCOMMITTED, READ COMMITTED, REPEATABLE READ, SERIALIZABLE

Transaction characteristics in SQL

- ❖ **SERIALIZABLE: Strict 2PL**
 - Ensures that T reads only changes made by committed transactions
 - No value **read or written** by T is changed by other transactions
 - Avoids a phantom phenomenon.
- ❖ **REPEATABLE READ: Does not do index locking.**
 - First and second are the same
 - But, T could experience a phantom phenomenon.
- ❖ **READ COMMITTED: Read locks are released immediately, exclusive locks are held till the end**
 - First is same
 - No value **written** by T is modified by other transactions
 - However, the value read by T may be modified by other transactions while T is in progress.
 - T could experience a phantom phenomenon.
- ❖ **READ UNCOMMITTED: does not obtain shared locks and does not request exclusive locks.**

Level	Dirty Read	Unrepeatable read	Phantom
READ UNCOMMITTED	Maybe	Maybe	Maybe
READ COMMITTED	No	Maybe	Maybe
REPEATABLE READ	No	No	Maybe
SERIALIZABLE	No	No	No

Transaction Characteristics in SQL

- ❖ **SERIALIZABLE ISOLATION** level is the safest and recommended for most transactions.
- ❖ For some transactions we can run at lower levels to improve the performance.
- ❖ Example: A statistical query which finds an average sailors age can be run at **READ COMMITTED** level or even **READ UNCOMMITTED** level
- ❖ The **ISOLATION LEVEL** and **ACCESS MODE** can be specified using the **SET TRANSACTION** command.
- ❖ Example: The following command declares the current transaction to be **SERIALIZABLE** and **READ ONLY**
SET TRANSACTION ISOLATION LEVEL SERIALIZABLE READ ONLY
- ❖ When transaction is started the default is **SERIALIZABLE** and **READ WRITE**

Introduction to Crash Recovery

- ❖ Recovery Manager is responsible for atomicity and durability.
- ❖ It ensures atomicity by undoing the actions of transactions
- ❖ It ensure durability by making sure that all actions of committed transactions survive system crashes and media failures.
- ❖ After the restart, the recovery manager is given the control to bring the database to a consistent state.
- ❖ Write ahead logging protocol is employed

Summary

- ❖ Concurrency control and recovery are among the most important functions provided by a DBMS.
- ❖ Users need not worry about concurrency.
 - System automatically inserts **lock/unlock** requests and schedules actions of different transactions to ensure that the resulting execution is **serializable**
 - **logging** is used to undo the actions of aborted transactions and to restore the system to a consistent state after a crash.