

Tree-Structured Indexes

Chapter 9

Outline

- ❖ Introduction
- ❖ Intuition for Tree Indexes
- ❖ Indexed sequential access method (IASM)
- ❖ B+ Trees: A dynamic Index Structure
- ❖ Search
- ❖ Insert
- ❖ Delete
- ❖ Duplicates
- ❖ B+ Trees in practice

Introduction

- ❖ *As for any index, 3 alternatives for data entries k^* :*
 - Data record with key value k
 - $\langle k, \text{rid of data record with search key value } k \rangle$
 - $\langle k, \text{list of rids of data records with search key } k \rangle$
- ❖ Choice is orthogonal to the *indexing technique* used to locate data entries k^* .
- ❖ Tree-structured indexing techniques support efficient insertion, deletion, *range searches* and *equality searches*.
- ❖ *ISAM (Indexed Sequential Access Method)*: static structure;
 - Efficient when the file is not frequently updated and unsuitable for files which grow and shrink.
- ❖ *B+ tree*: dynamic, adjusts gracefully under inserts and deletes.
- ❖ In both structures, leaf pages contain data entries.

Intuition

- ❖ *“Find all students with $gpa > 3.0$ ”*
 - If data is in sorted file, do binary search to find first such student, then scan to find others.
 - Cost of binary search can be quite high.
 - Cost is proportional to number of pages fetched.
- ❖ Simple idea: Create an ‘index’ file with one record per page in the original data file, of the form <first key on the page, pointer to page>




👉 *Can do binary search on (smaller) index file!*


Intuition...

- ❖ Index file is much smaller than the data file
- ❖ Binary search of index file is much faster than the binary search of data file.
- ❖ However, binary search of index file is still expensive.
- ❖ Large index file motivates the tree index idea.
 - Applying the idea of auxiliary structure on the collection of index records, until smallest structure is in one page.
 - It results into a tree!
- ❖ Locating a record involves a traversal from the root to leaf, with one I/O per level.
- ❖ Root will be in the buffer pool
- ❖ Given a typical fan-out value (over 100), trees rarely have more than 3-4 levels.

ISAM



❖ Index file may still be quite large. But we can apply the idea repeatedly!



➡ Leaf pages contain *data entries*.

Comments on ISAM

Data
Pages


Index Pages

Overflow pages

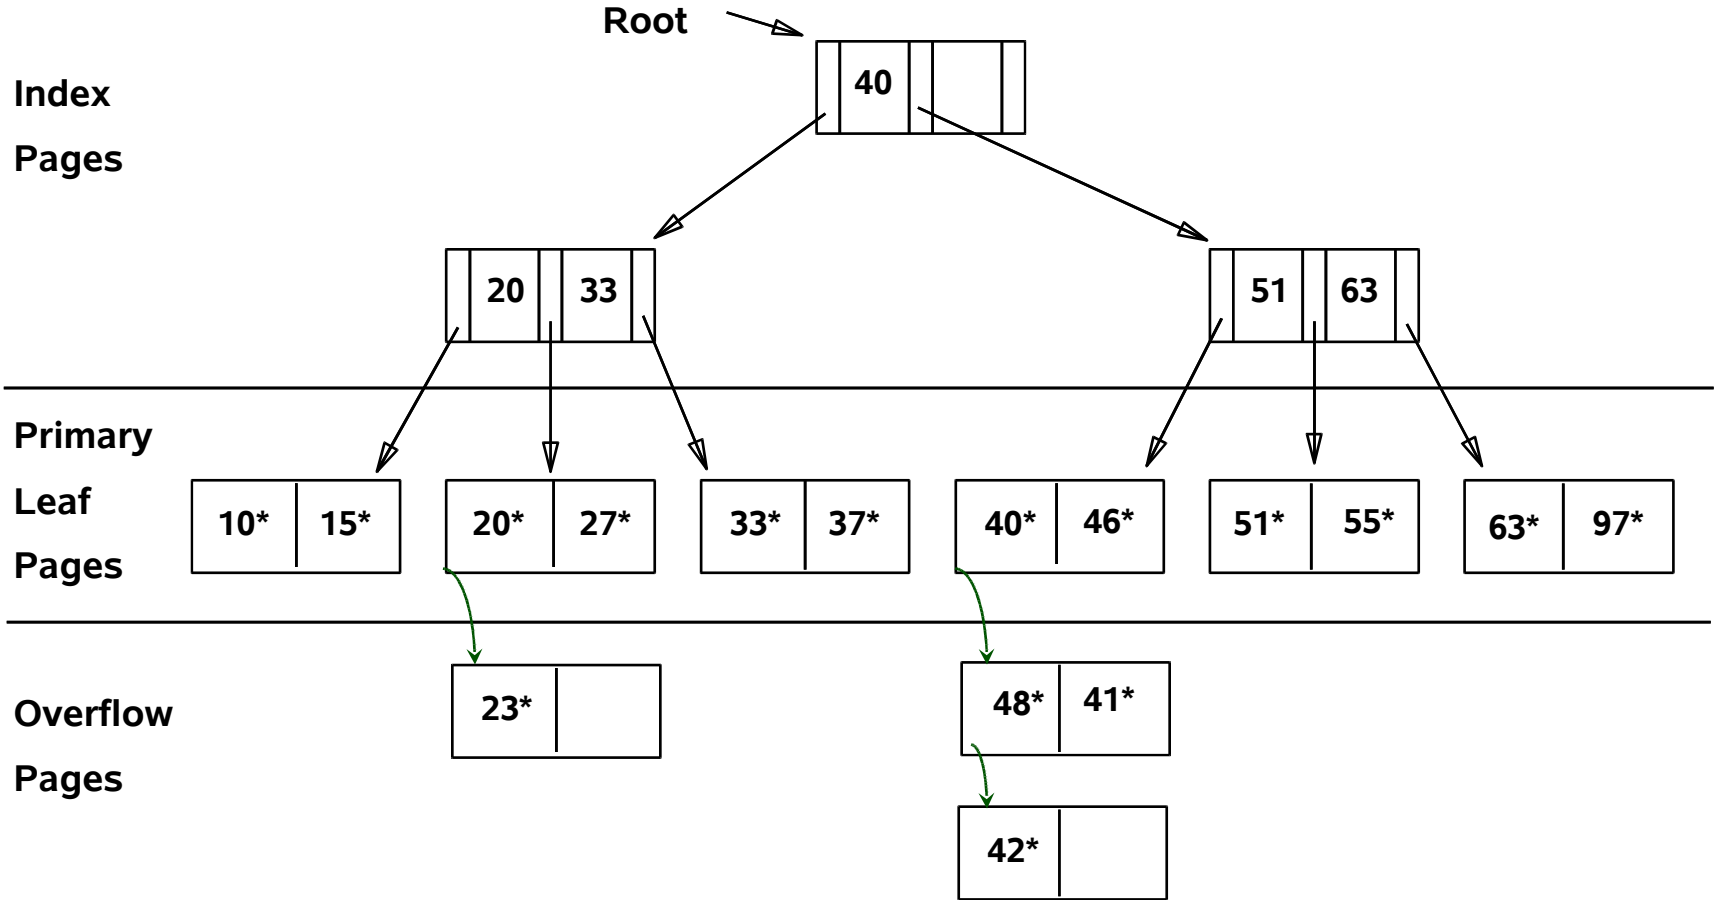
- ❖ *File creation*: Leaf (data) pages allocated sequentially, sorted by search key; then index pages allocated, then space for overflow pages.
 - ❖ *Index entries*: **<search key value, page id>**; they **`direct' search for data entries**, which are in leaf pages.
 - ❖ *Search*: Start at root; use key comparisons to go to leaf. Cost $\log_F N$; $F = \# \text{ entries/index pg}$, $N = \# \text{ leaf pgs}$
 - ❖ *Insert*: Find leaf data entry belongs to, and put it there.
 - ❖ *Delete*: Find and remove from leaf; if empty overflow page, de-allocate.
- ❖ **Static tree structure**: *inserts/deletes affect only leaf pages.*

Example ISAM Tree

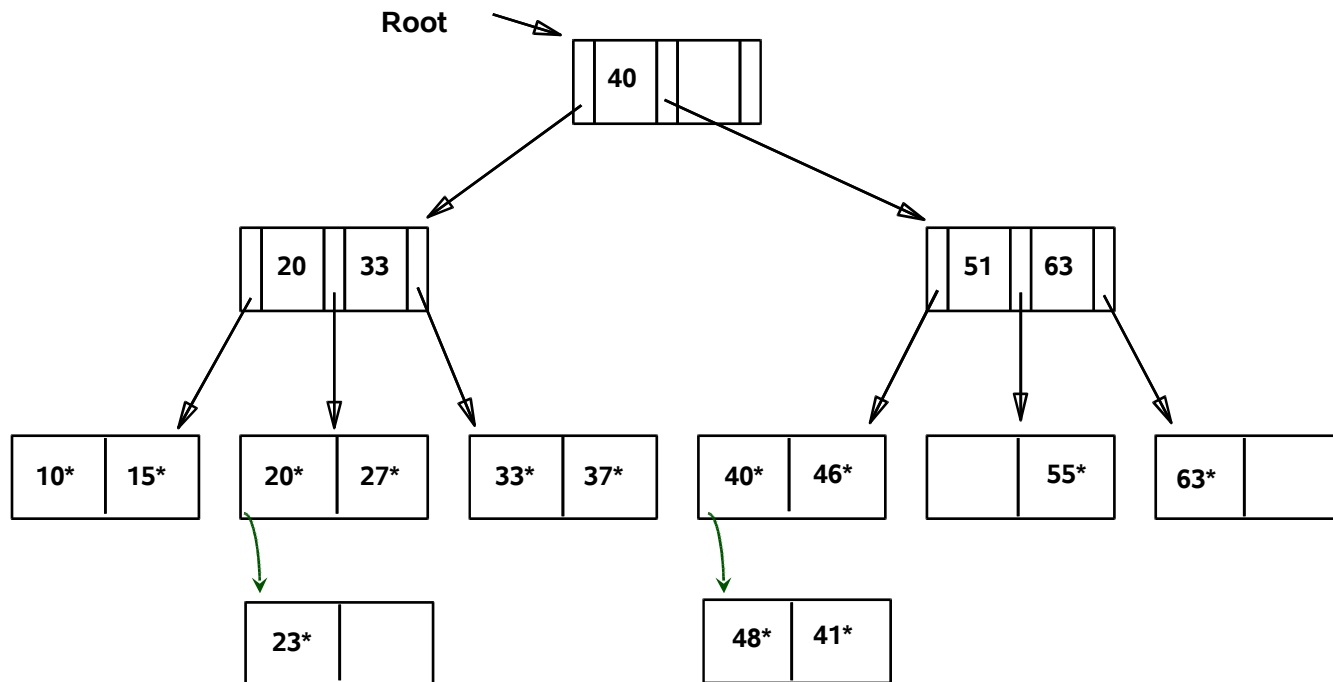
- ❖ Each node can hold 2 entries; no need for 'next-leaf-page' pointers. (Why?)



After Inserting 23*, 48*, 41*, 42* ...



... Then Deleting 42*, 51*, 97*




➡ Note that 51* appears in index levels, but not in leaf!

ISAM--Discussion

- ❖ More inserts and deletes results into long overflow chains, which increases the retrieval time because overflow chains have to searched.
- ❖ In case of concurrent searches, it offers significant advantages.
 - Need not lock index pages as they are not modified.
 - If the data distribution size is static with no overflow chains, ISAM may be preferable to B+ trees.

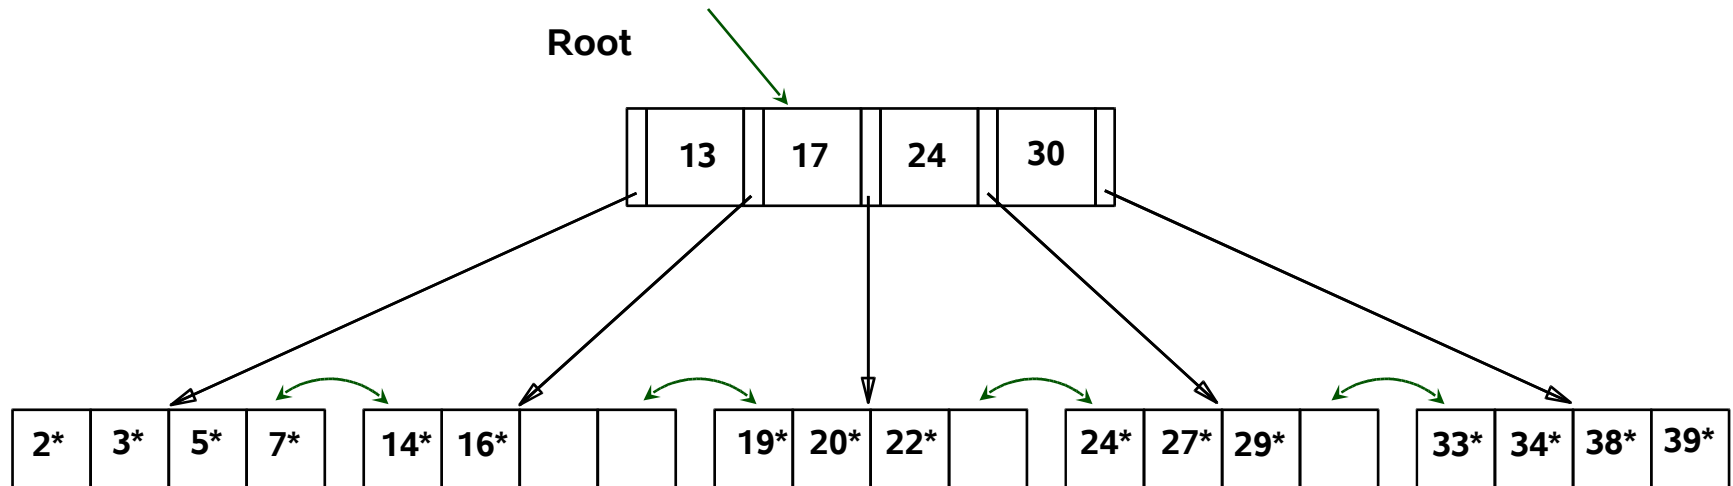
B+ Tree: Most Widely Used Index

- ❖ Insert/delete at $\log_F N$ cost; keep tree *height-balanced*. (F = fanout, N = # leaf pages)
- ❖ Minimum 50% occupancy (except for root). Each node contains $\mathbf{d} \leq \underline{m} \leq 2\mathbf{d}$ entries. The parameter \mathbf{d} is called the *order* of the tree.
- ❖ Supports equality and range-searches efficiently.



Example B+ Tree

- ❖ Search begins at root, and key comparisons direct it to a leaf (as in ISAM).
- ❖ Search for 5*, 15*, all data entries $\geq 24^*$...



👉 *Based on the search for 15*, we know it is not in the tree!*

B+ Trees in Practice



- ❖ Typical order: 100. Typical fill-factor: 67%.
 - average fanout = 133
- ❖ Typical capacities:
 - Height 4: $133^4 = 312,900,700$ records
 - Height 3: $133^3 = 2,352,637$ records
- ❖ Can often hold top levels in buffer pool:
 - Level 1 = 1 page = 8 Kbytes
 - Level 2 = 133 pages = 1 Mbyte
 - Level 3 = 17,689 pages = 133 MBytes

Inserting a Data Entry into a B+ Tree


- ❖ Find correct leaf L .
- ❖ Put data entry onto L .
 - If L has enough space, *done!*
 - Else, must split L (into L and a new node $L2$)
 - Redistribute entries evenly, copy up middle key.
 - Insert index entry pointing to $L2$ into parent of L .
- ❖ This can happen recursively
 - To split index node, redistribute entries evenly, but push up middle key. (Contrast with leaf splits.)
- ❖ Splits “grow” tree; root split increases height.
 - Tree growth: gets wider or one level taller at top.

Inserting 8* into Example B+ Tree

- ❖ Observe how minimum occupancy is guaranteed in both leaf and index pg splits.
- ❖ Note difference between *copy-up* and *push-up*; be sure you understand the reasons for this.



Example B+ Tree After Inserting 8*




Notice that root was split, leading to increase in height.

In this example, we can avoid split by re-distributing entries; however, this is usually not done in practice.

Deleting a Data Entry from a B+ Tree

- ❖ Start at root, find leaf L where entry belongs.
- ❖ Remove the entry.
 - If L is at least half-full, *done!*
 - If L has only **d-1** entries,
 - Try to **re-distribute**, borrowing from sibling (*adjacent node with same parent as L*).
 - If re-distribution fails, **merge** L and sibling.
- ❖ If merge occurred, must delete entry (pointing to L or sibling) from parent of L .
- ❖ Merge could propagate to root, decreasing height.



Example Tree After (Inserting 8*, Then) Deleting 19* and 20* ...



- ❖ Deleting 19* is easy.
- ❖ Deleting 20* is done with re-distribution. Notice how middle key is *copied up*.


... And Then Deleting 24*

- ❖ Must merge.
- ❖ Observe *'toss'* of index entry (on right), and *'pull down'* of index entry (below).




Example of Non-leaf Re-distribution

- ❖ Tree is shown below *during deletion* of 24*. (What could be a possible initial tree?)
- ❖ In contrast to previous example, can re-distribute entry from left child of root to right child.



After Re-distribution

- ❖ Intuitively, entries are **re-distributed by 'pushing through'** the splitting entry in the parent node.
- ❖ It suffices to re-distribute index entry with key 20; we've re-distributed 17 as well for illustration.




Prefix Key Compression

- ❖ Important to increase fan-out. (Why?)
- ❖ Key values in index entries only `direct traffic`; can often compress them.
 - E.g., If we have adjacent index entries with search key values *Dannon Yogurt*, *David Smith* and *Devarakonda Murthy*, we can abbreviate *David Smith* to *Dav*. (The other keys can be compressed too ...)
 - Is this correct? Not quite! What if there is a data entry *Davey Jones*? (Can only compress *David Smith* to *Davi*)
 - In general, while compressing, must leave each index entry greater than every key value (in any subtree) to its left.
- ❖ Insert/delete must be suitably modified.


Bulk Loading of a B+ Tree

- ❖ If we have a large collection of records, and we want to create a B+ tree on some field, doing so by repeatedly inserting records is very slow.
- ❖ Bulk Loading can be done much more efficiently.
- ❖ *Initialization*: Sort all data entries, insert pointer to first (leaf) page in a new (root) page.



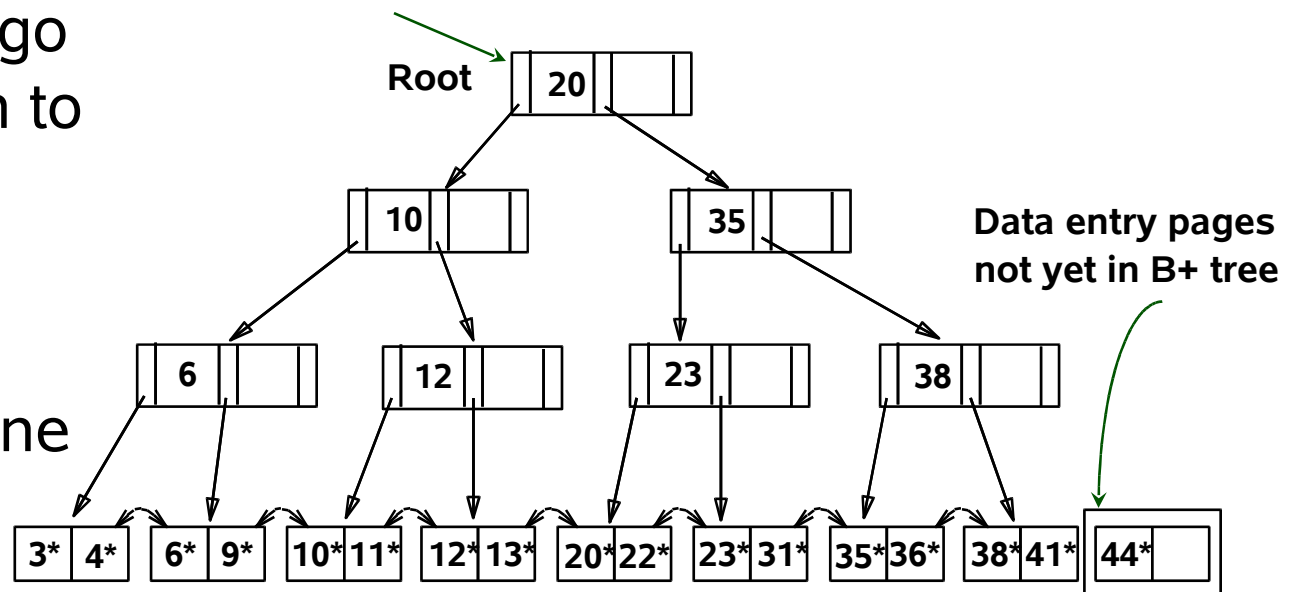
Bulk Loading (Contd.)

- ❖ Index entries for leaf pages always entered into right-most index page just above leaf level.



When this fills up, it splits. (Split may go up right-most path to the root.)

- ❖ Much faster than repeated inserts, especially when one considers locking!



Summary of Bulk Loading

- ❖ Option 1: multiple inserts.
 - Slow.
 - Does not give sequential storage of leaves.
- ❖ Option 2: Bulk Loading
 - Has advantages for concurrency control.
 - Fewer I/Os during build.
 - Leaves will be stored sequentially (and linked, of course).
 - Can control “fill factor” on pages.

A Note on `Order`

- ❖ *Order* (**d**) concept replaced by physical space criterion in practice (*`at least half-full`*).
 - Index pages can typically hold many more entries than leaf pages.
 - Variable sized records and search keys mean different nodes will contain different numbers of entries.
 - Even with fixed length fields, multiple records with the same search key value (*duplicates*) can lead to variable-sized data entries (if we use Alternative (3)).

Summary

- ❖ Tree-structured indexes are ideal for range-searches, also good for equality searches.
- ❖ ISAM is a static structure.
 - Only leaf pages modified; overflow pages needed.
 - Overflow chains can degrade performance unless size of data set and data distribution stay constant.
- ❖ B+ tree is a dynamic structure.
 - Inserts/deletes leave tree height-balanced; $\log_F N$ cost.
 - High fanout (**F**) means depth rarely more than 3 or 4.
 - Almost always better than maintaining a sorted file.

Summary (Contd.)

- Typically, 67% occupancy on average.
 - Usually preferable to ISAM, modulo *locking* considerations; adjusts to growth gracefully.
 - If data entries are data records, splits can change rids!
- ❖ Key compression increases fanout, reduces height.
 - ❖ Bulk loading can be much faster than repeated inserts for creating a B+ tree on a large data set.
 - ❖ Most widely used index in database management systems because of its versatility. One of the most optimized components of a DBMS.